

The Essence of Structured Streaming Computation

Joseph W. Cutler

Christopher Watson

Philip Hilliard

Harrison Goldstein

Caleb Stanford (UC Davis)

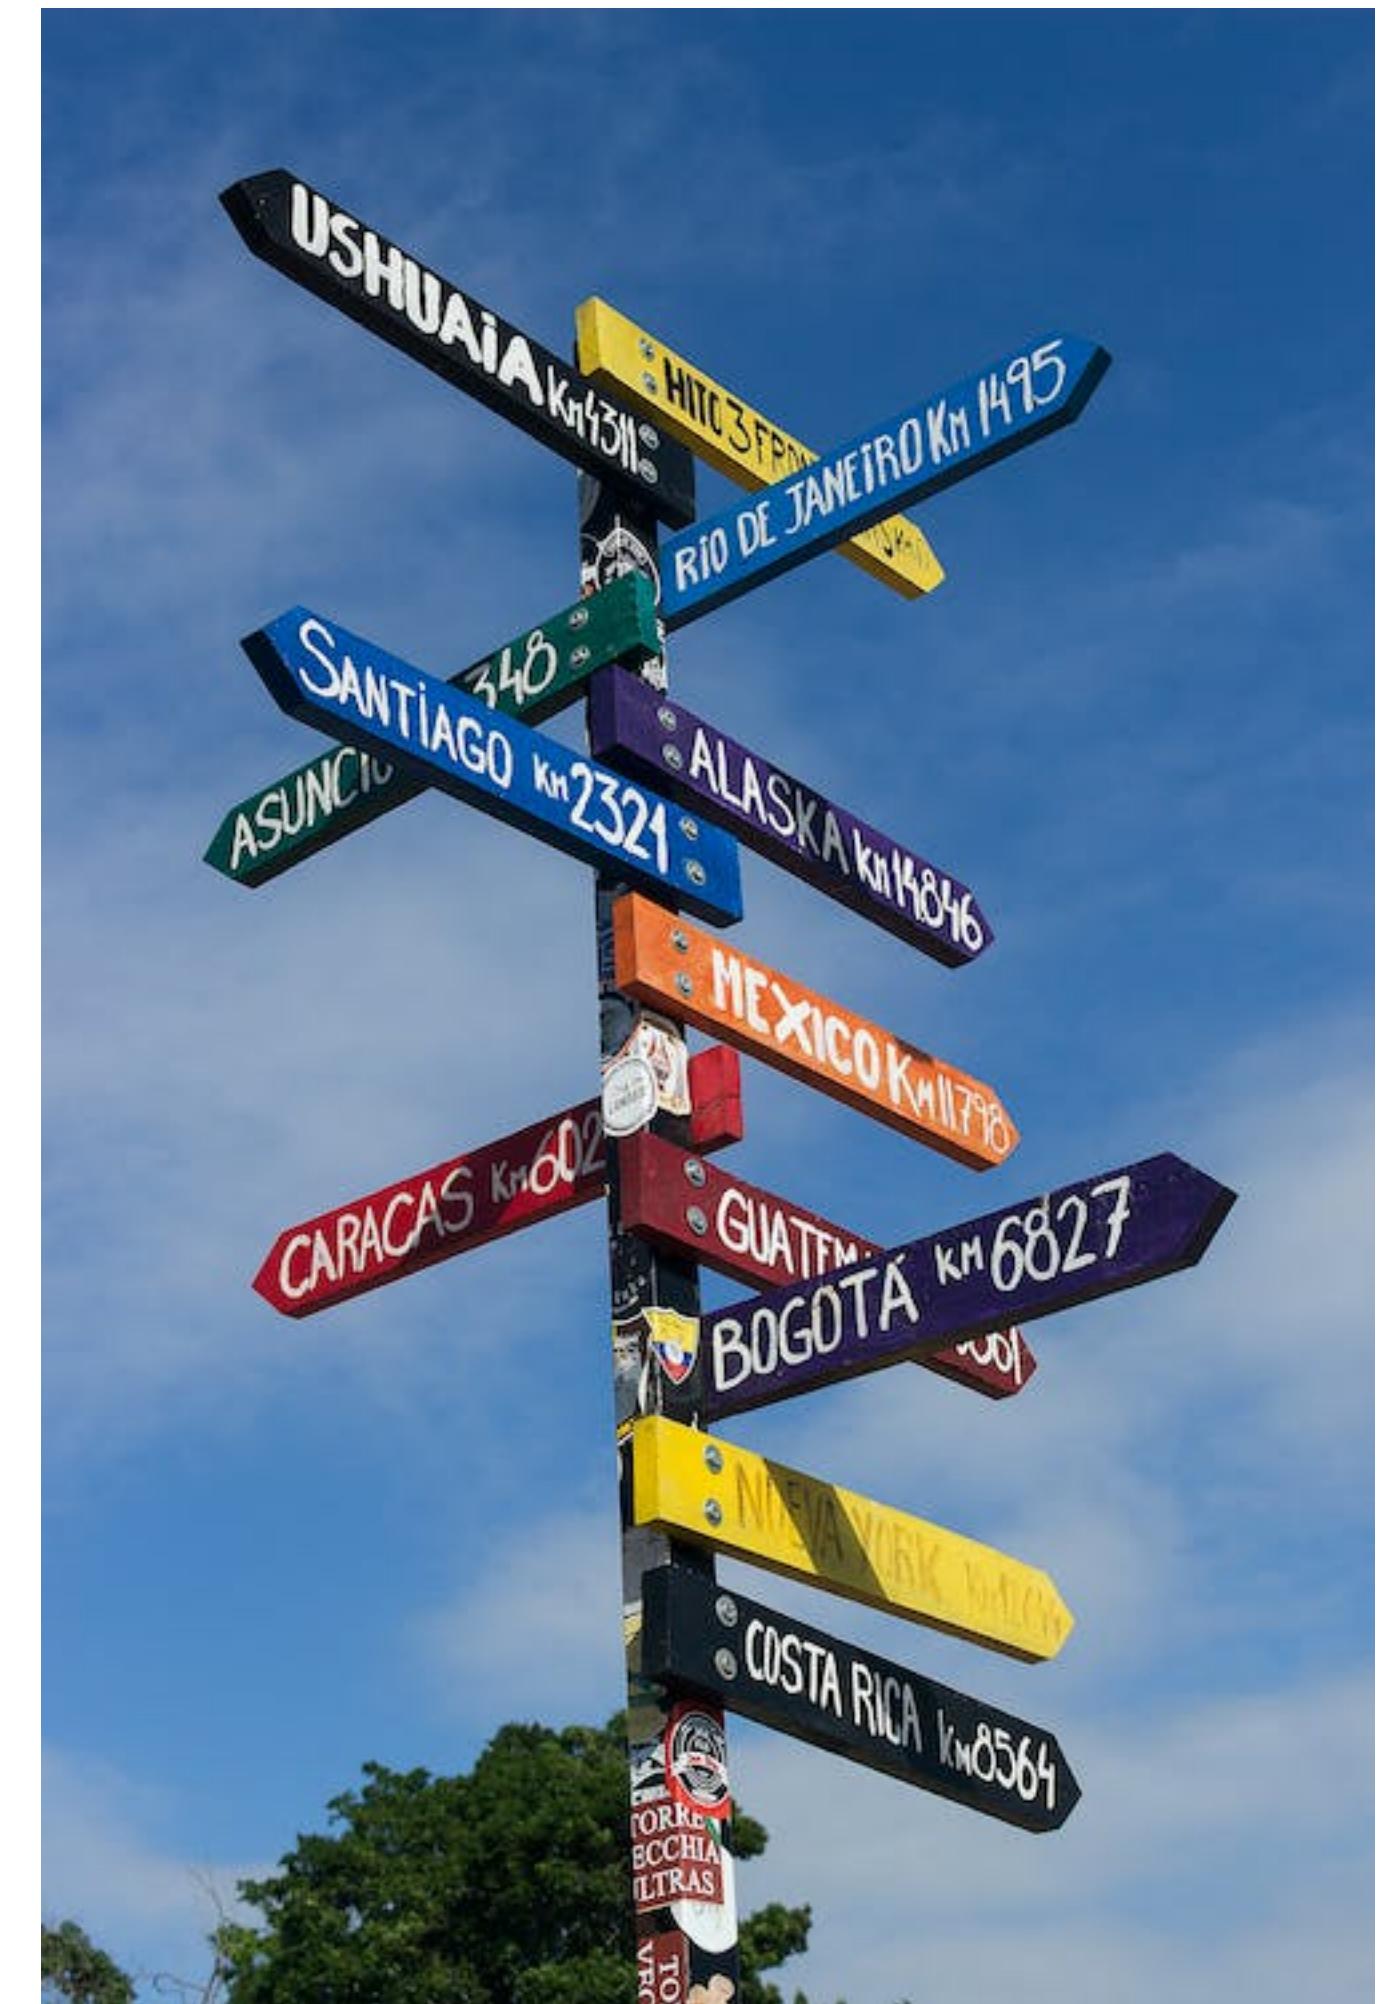
Benjamin C. Pierce

**Types are the essence of
Structured Streaming**

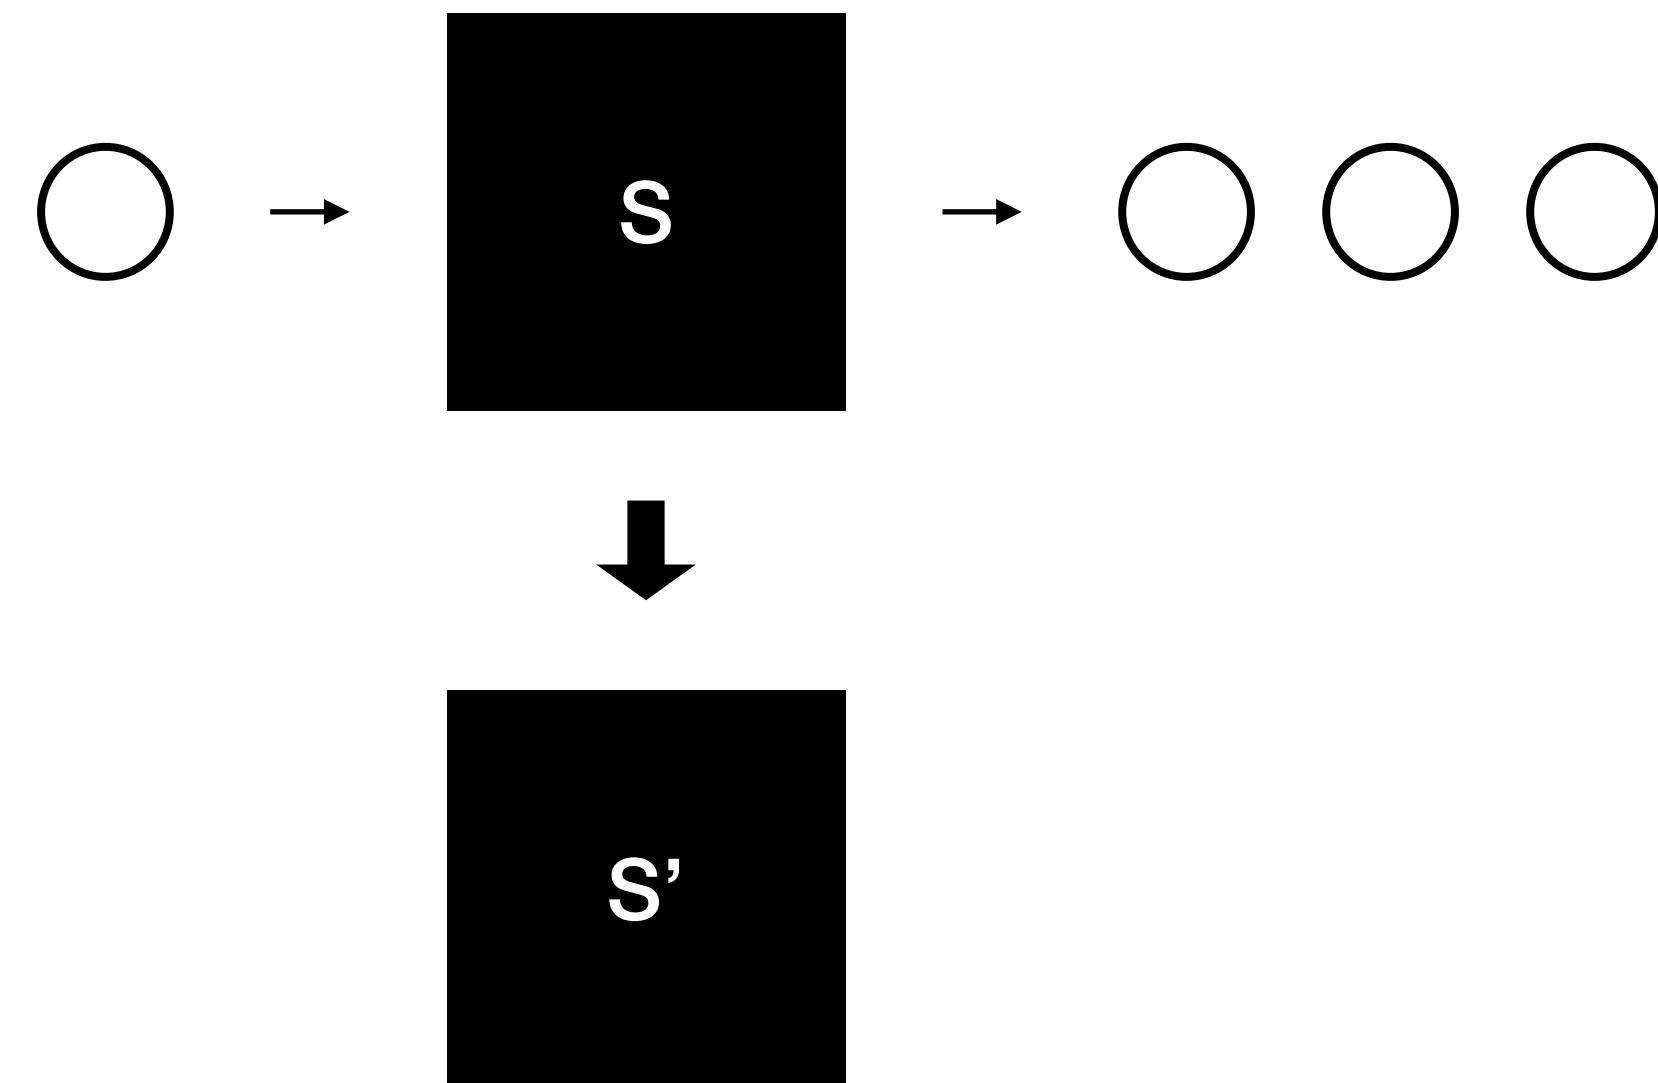
**Picking the right types for streams leads
you directly to idiomatic stream
programming**

Let's define some terms.

1. What is Streaming Computation?
2. What is this thing I'm calling *structured streaming*?
3. Our Types For Streams
4. The Stream-Typed Core Calculus

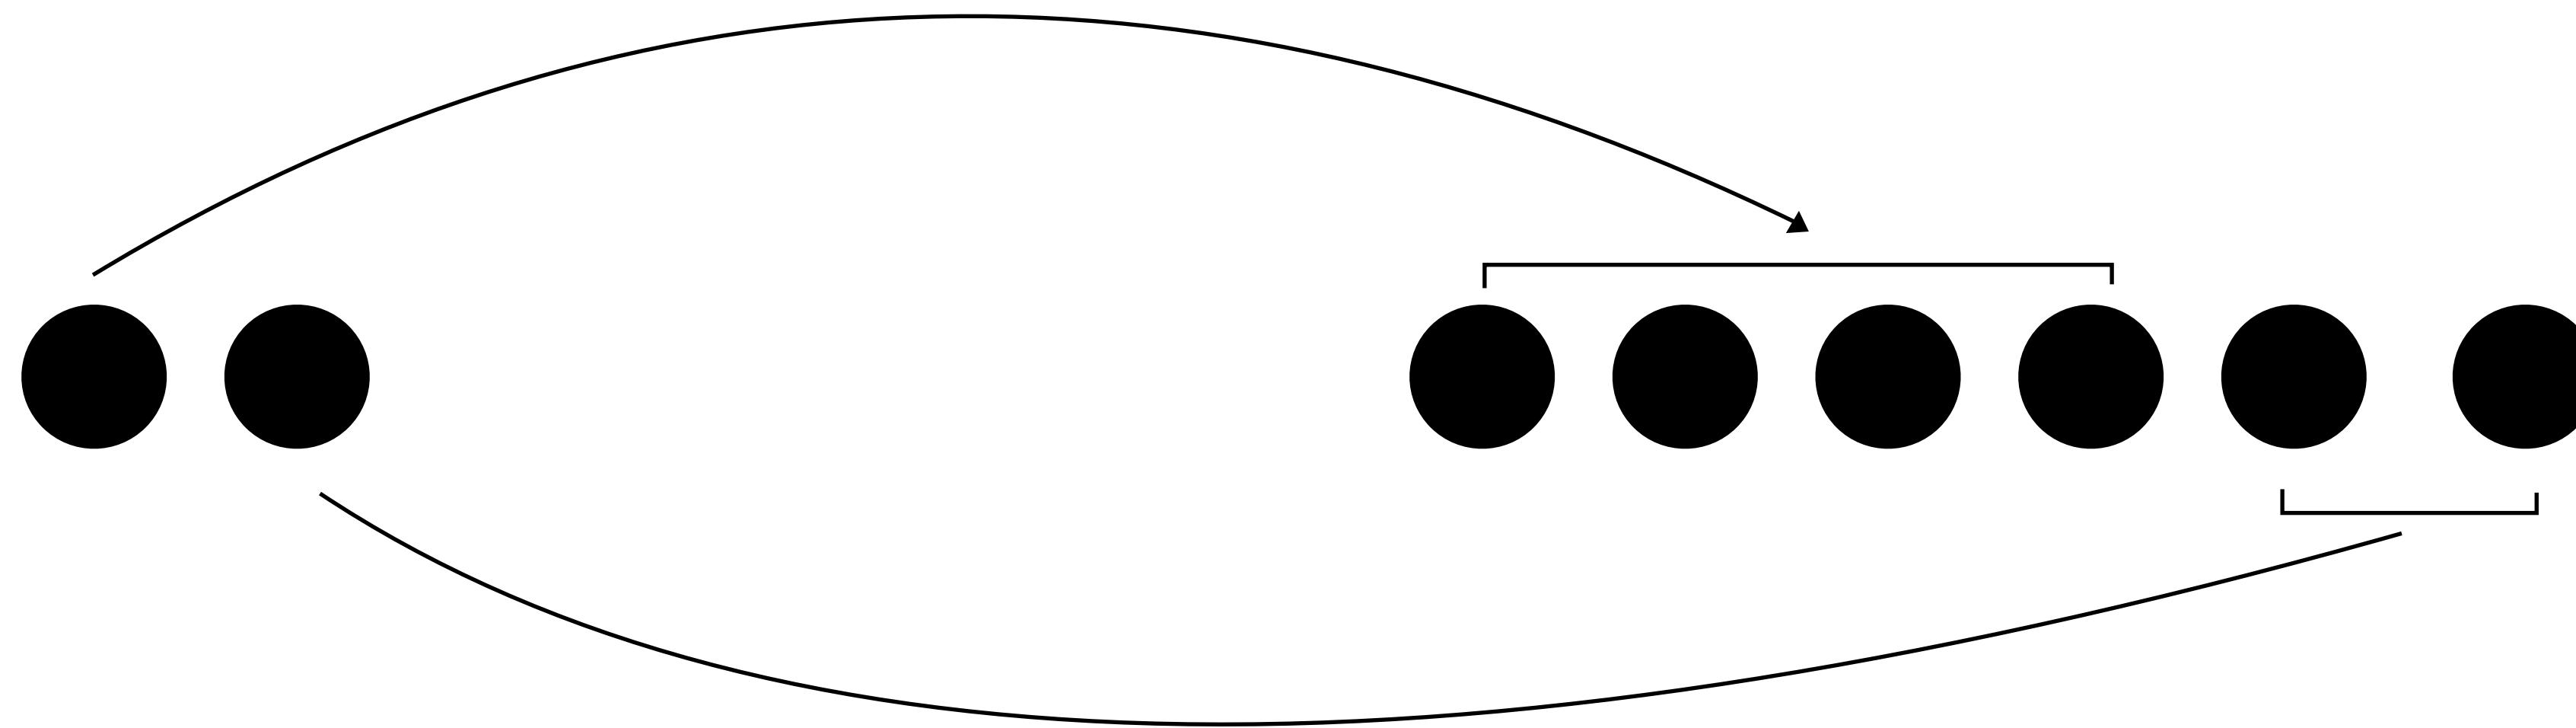


What is Streaming Computation?



$f : \text{state} \rightarrow \text{input} \rightarrow \text{state} * (\text{list output})$

Programming Model of Streaming Computation



Transform streams, event by event

Example: Streaming Data Analytics

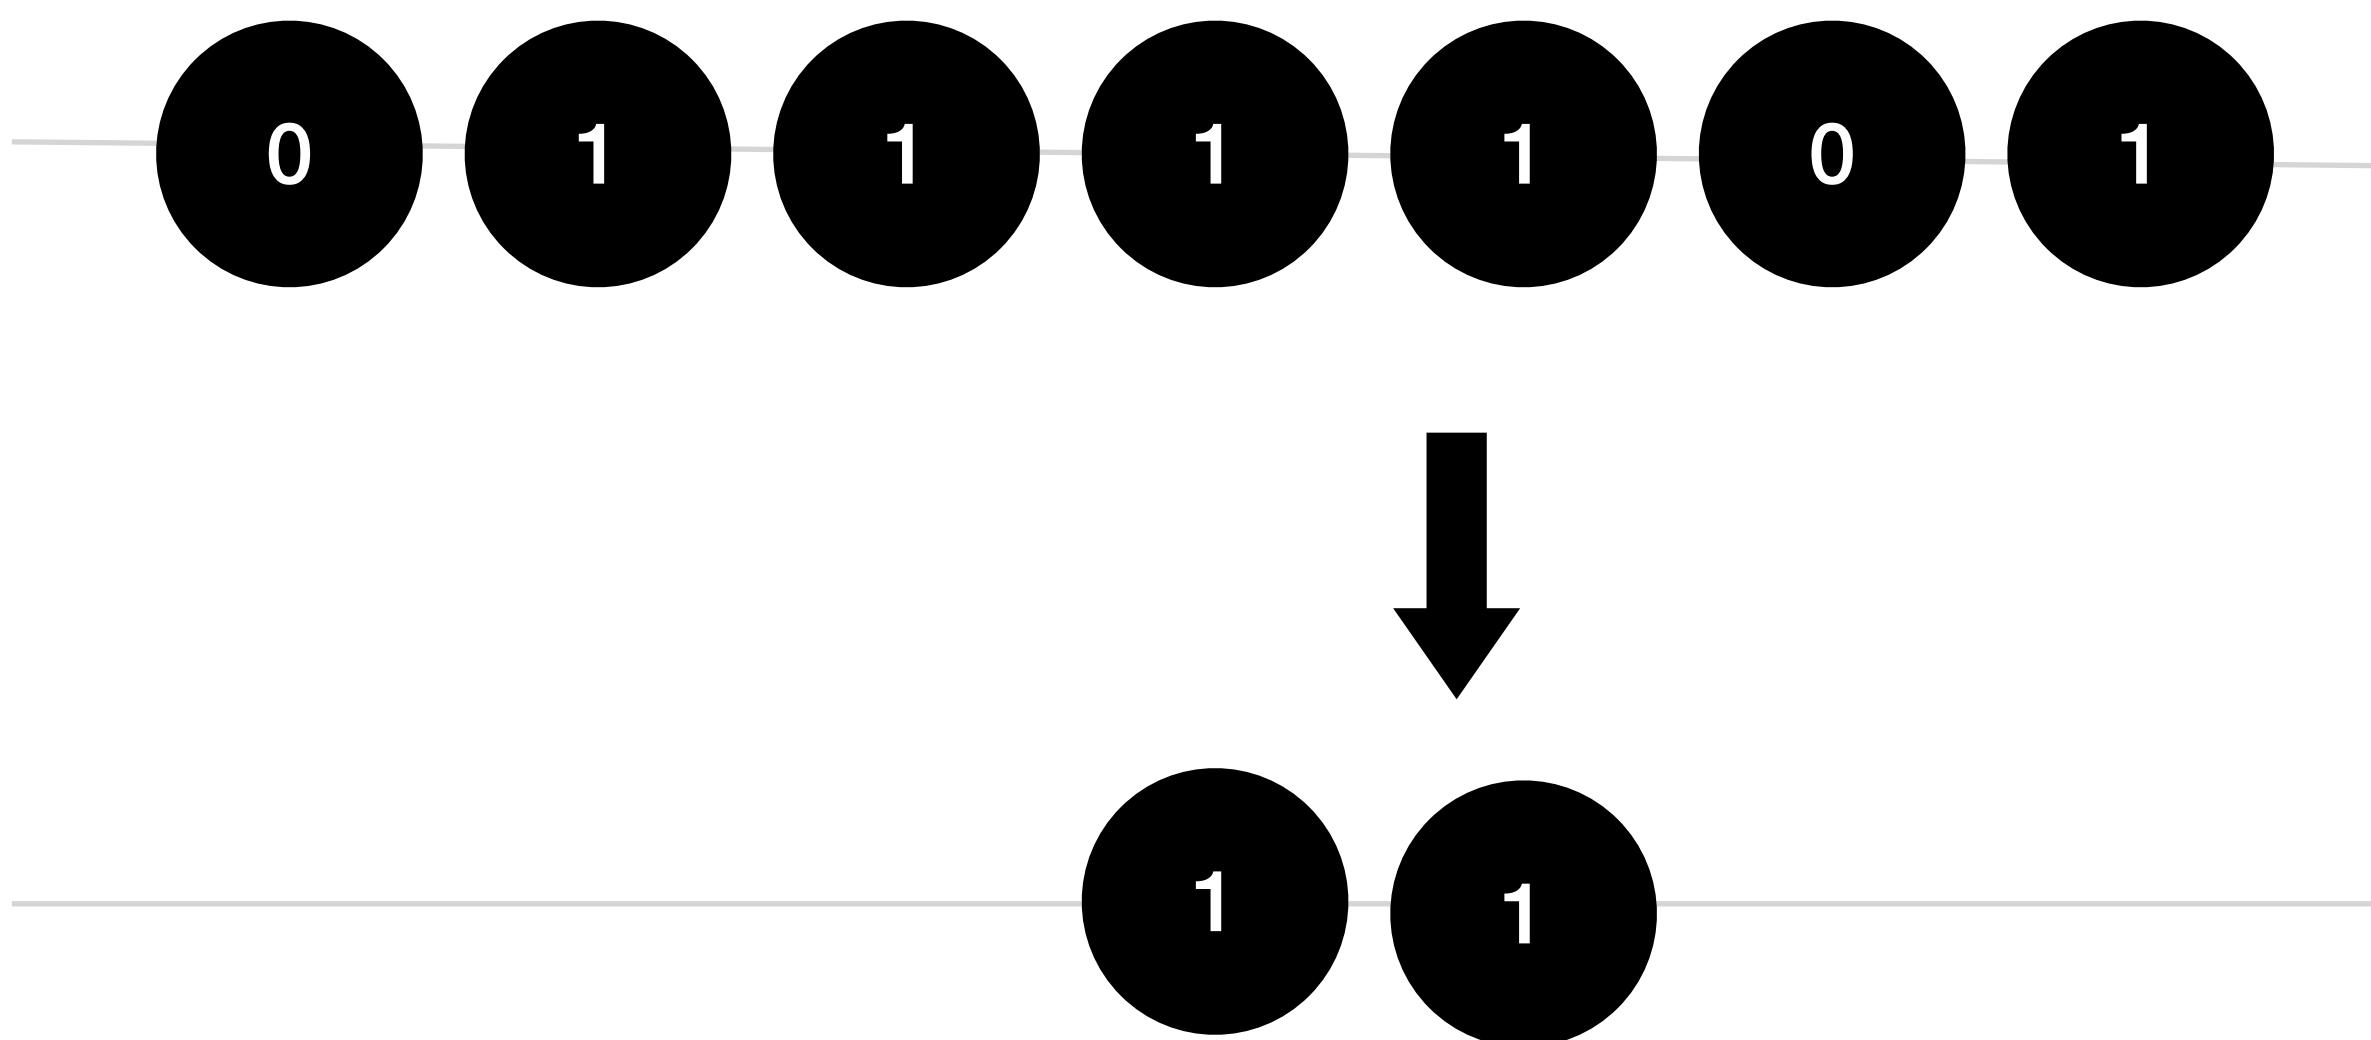


Example: Streaming Data Analytics

Motion Sensor



Alert with 1 if motion is detected in three subsequent periods



```
handle(bool x) {  
    if(x && prev1 && prev2) {  
        out(1);  
    }  
    prev2 = prev1;  
    prev1 = x;  
}
```

Manually-Written State Machine!

Example: ML Model Training

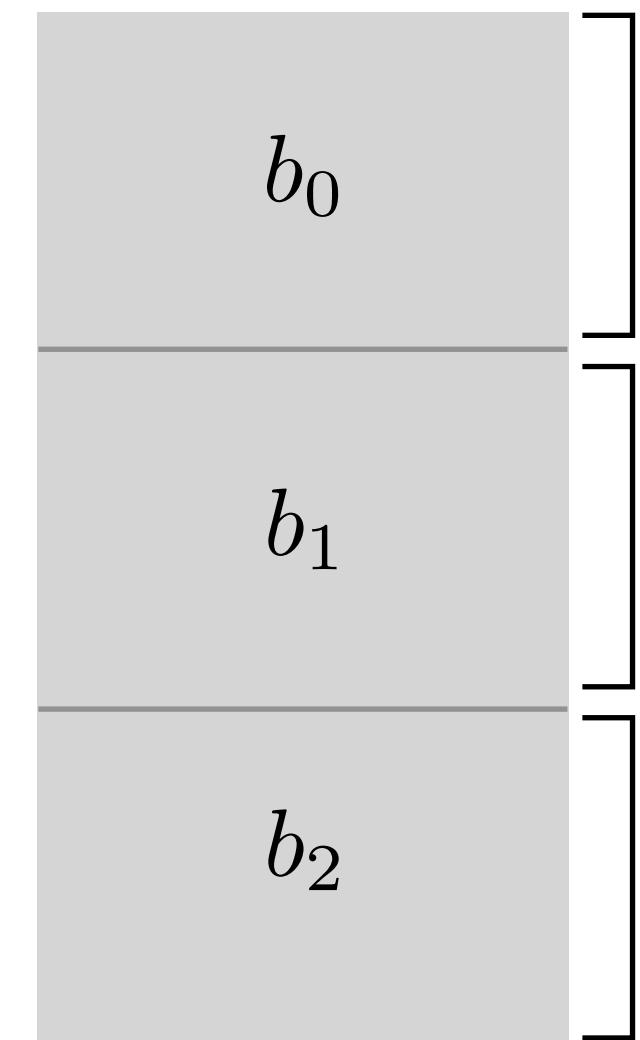


Example: ML Model Training

Minibatch Gradient Descent

Another Manually-Written State Machine!

```
handle(vec x) {  
    loss += gradloss(x, w)  
    n += 1  
    if (n == k) {  
        w -= (alpha / k) * loss  
        loss, n = 0  
    }  
}
```



$$w_1 = w_0 - \frac{\alpha}{k} \sum_{x \in b_0} \nabla L(x, w_0)$$

$$w_2 = w_1 - \frac{\alpha}{k} \sum_{x \in b_1} \nabla L(x, w_1)$$

$$w_3 = w_2 - \frac{\alpha}{k} \sum_{x \in b_2} \nabla L(x, w_2)$$

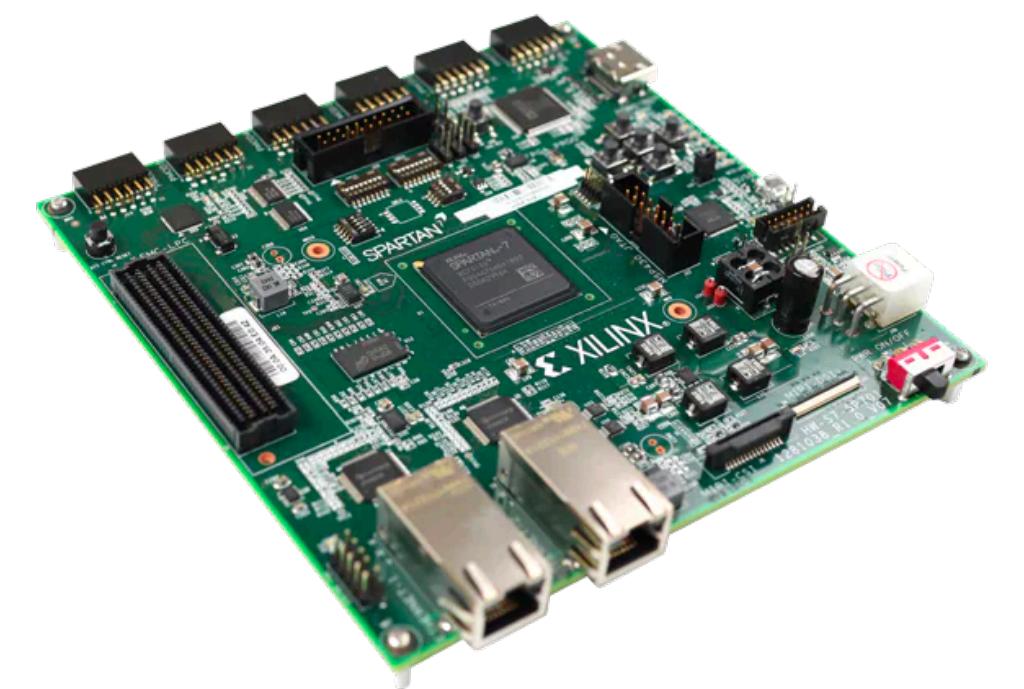
1. Accumulate the loss
2. Update the weights when we've seen k vectors

Example: Networking Hardware



Example: Networking Hardware

01000101010110... →



→ 1100111101001...

*Extremely Manually-
written state
machine!*

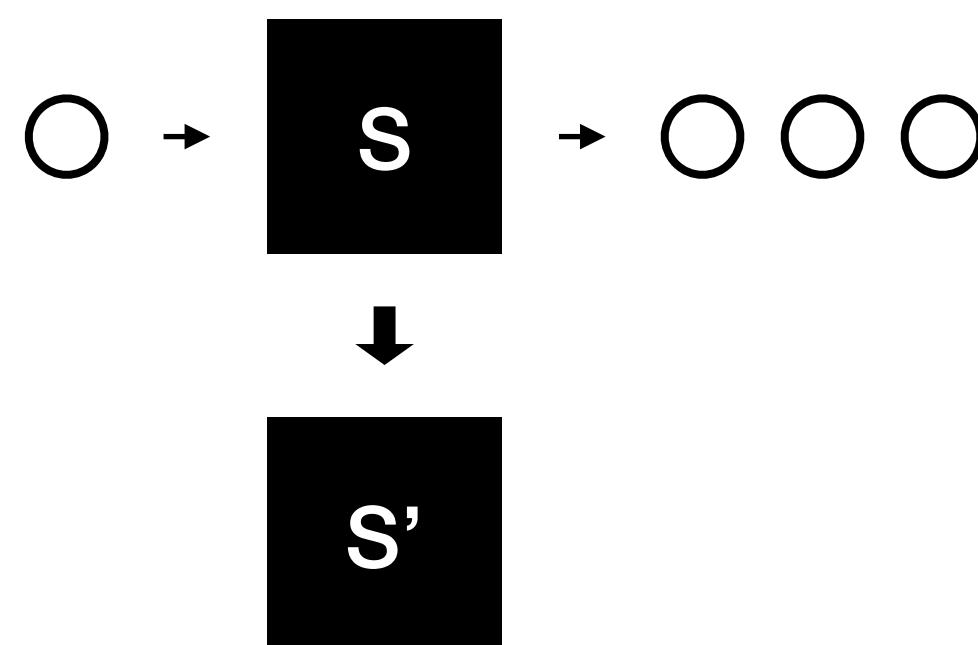
**In Summary: Powerful but
Challenging!**

What is *Structured Streaming* Computation?

Richer Programming Model

stream a \rightarrow stream b

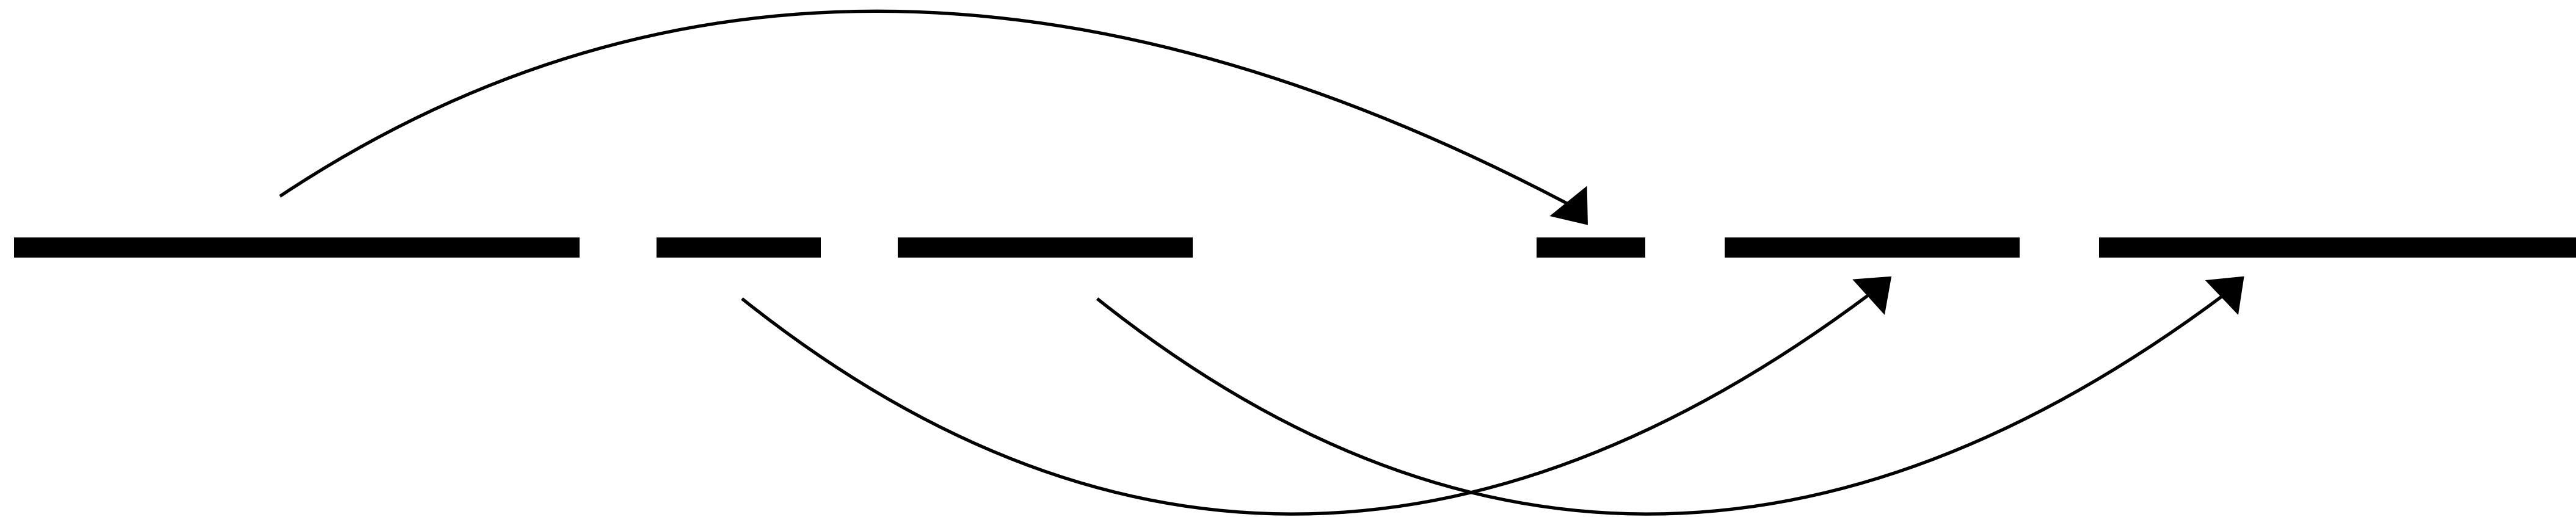
map, filter, reduce, *windows*, and *structured types*



Computational model remains the same

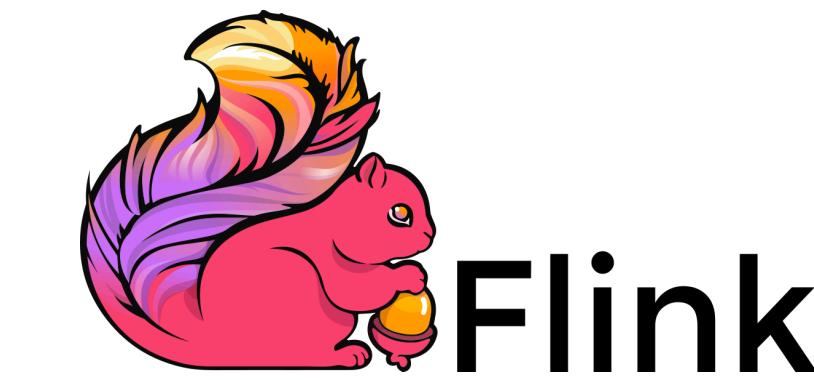
Programming Model of Structured Streaming

Transforming whole extents of a stream at a time



Extents themselves may have variable size, or be composed of structured types

Example: Structured Data Analytics



map : $(a \rightarrow b) \rightarrow \text{stream } a \rightarrow \text{stream } b$

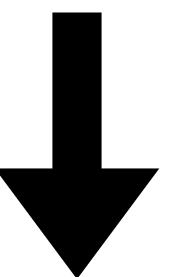
reduce : $(a \rightarrow b \rightarrow b) \rightarrow b \rightarrow \text{stream } a \rightarrow \text{stream } b$

slidingWindow : $\text{int} \rightarrow (\text{list } a \rightarrow b) \rightarrow \text{stream } a \rightarrow \text{stream } b$

Example: Structured Streaming Data Analytics

Motion Sensor

```
handle(bool x) {  
    if(x && prev1 && prev2) {  
        out(1);  
    }  
    prev2 = prev1;  
    prev1 = x;  
}
```

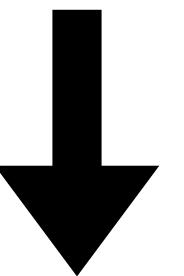


```
input.slidingWindow(3, all)
```

Example: Structured ML Model Training

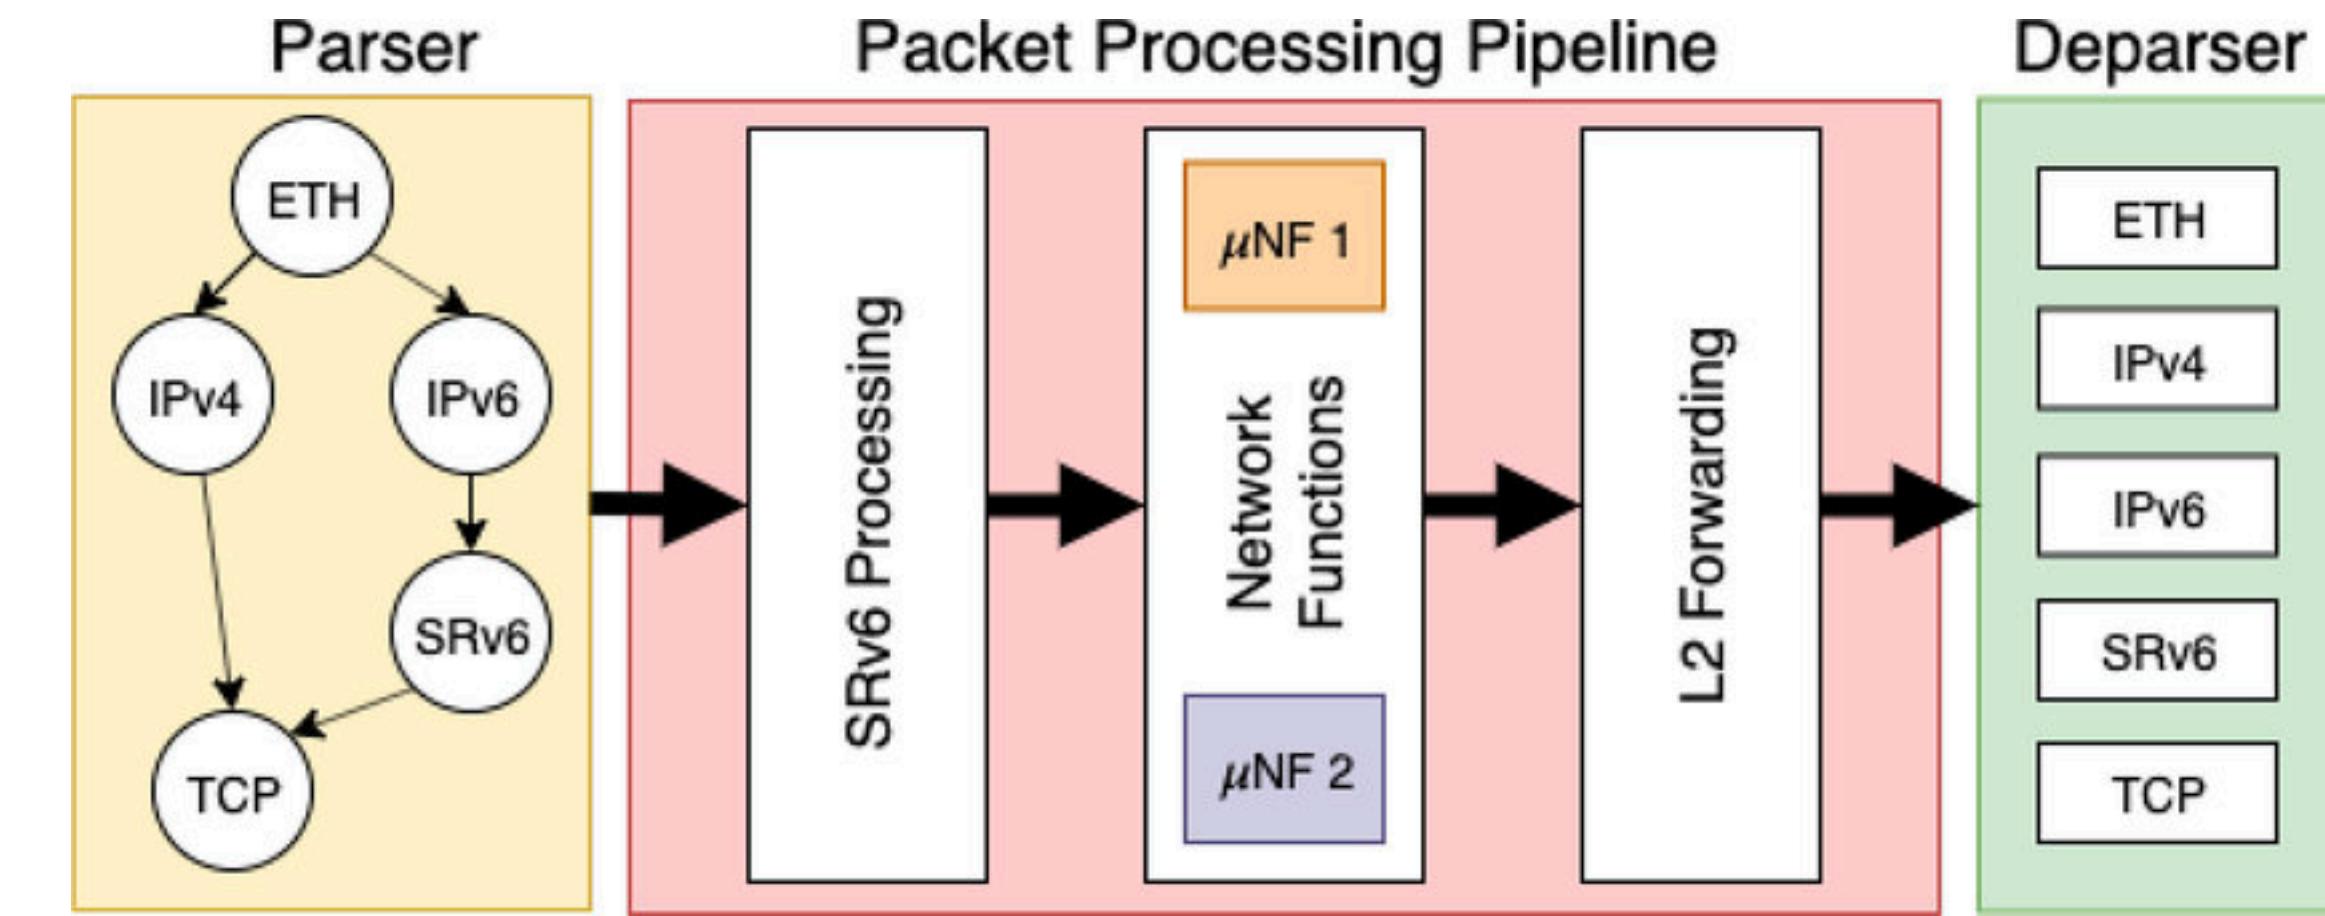
Minibatch Gradient Descent

```
handle(vec x) {  
    loss += gradloss(x, w)  
    n += 1  
    if (n == k) {  
        w -= alpha * (loss / k)  
        loss, n = 0  
    }  
}
```



```
training.map(gradloss).tumblingWindow(k, average_and_update)
```

Example: Structured Networking Hardware Streaming



parse : stream bit \rightarrow stream IPv4

process : stream IPv4 \rightarrow stream IPv4

deparse : stream IPv4 \rightarrow stream bit

Structured Streaming

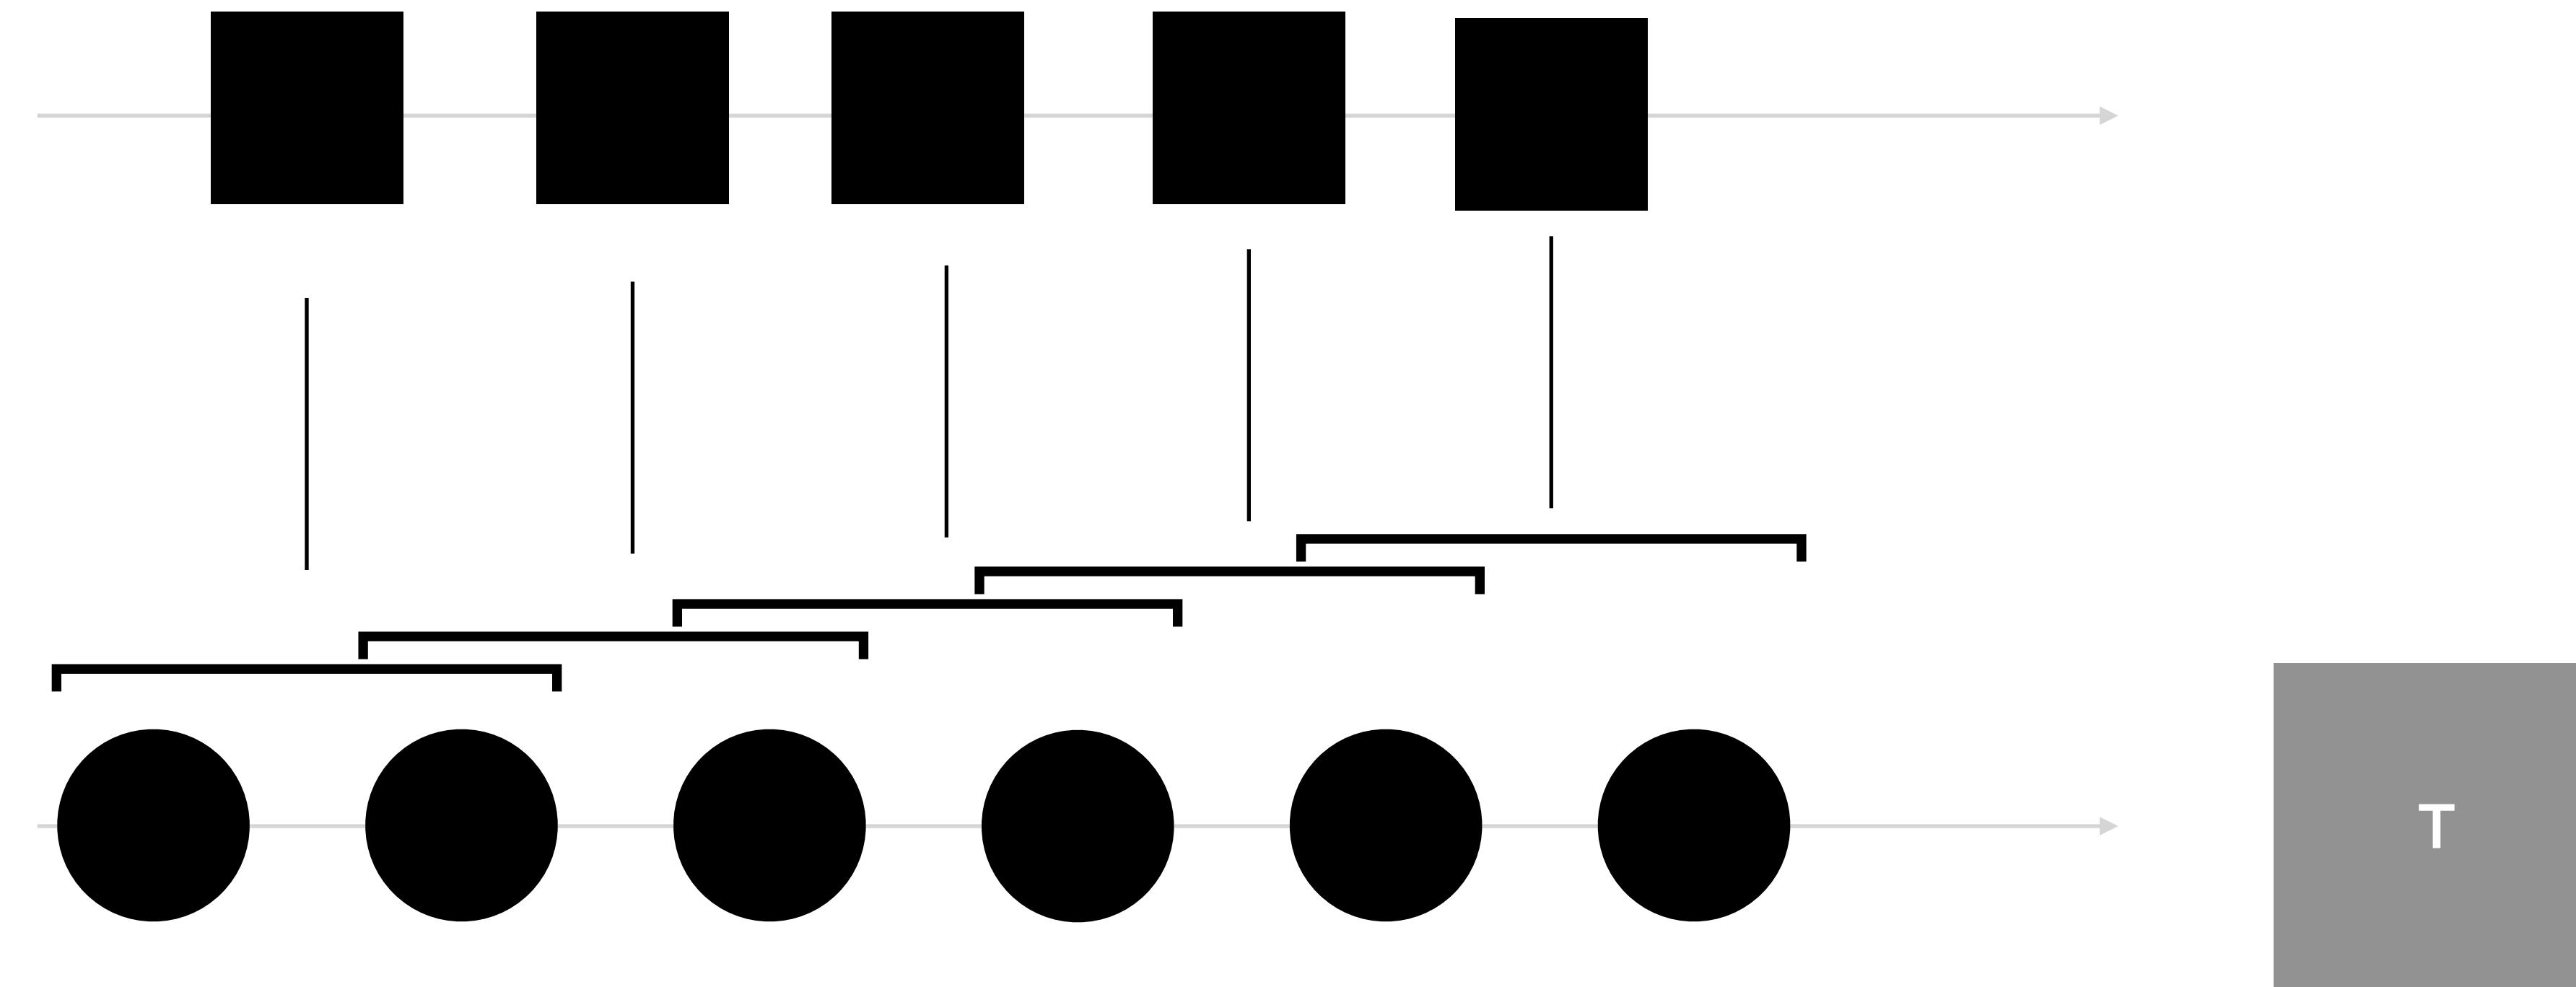
requires

...

Structured Types

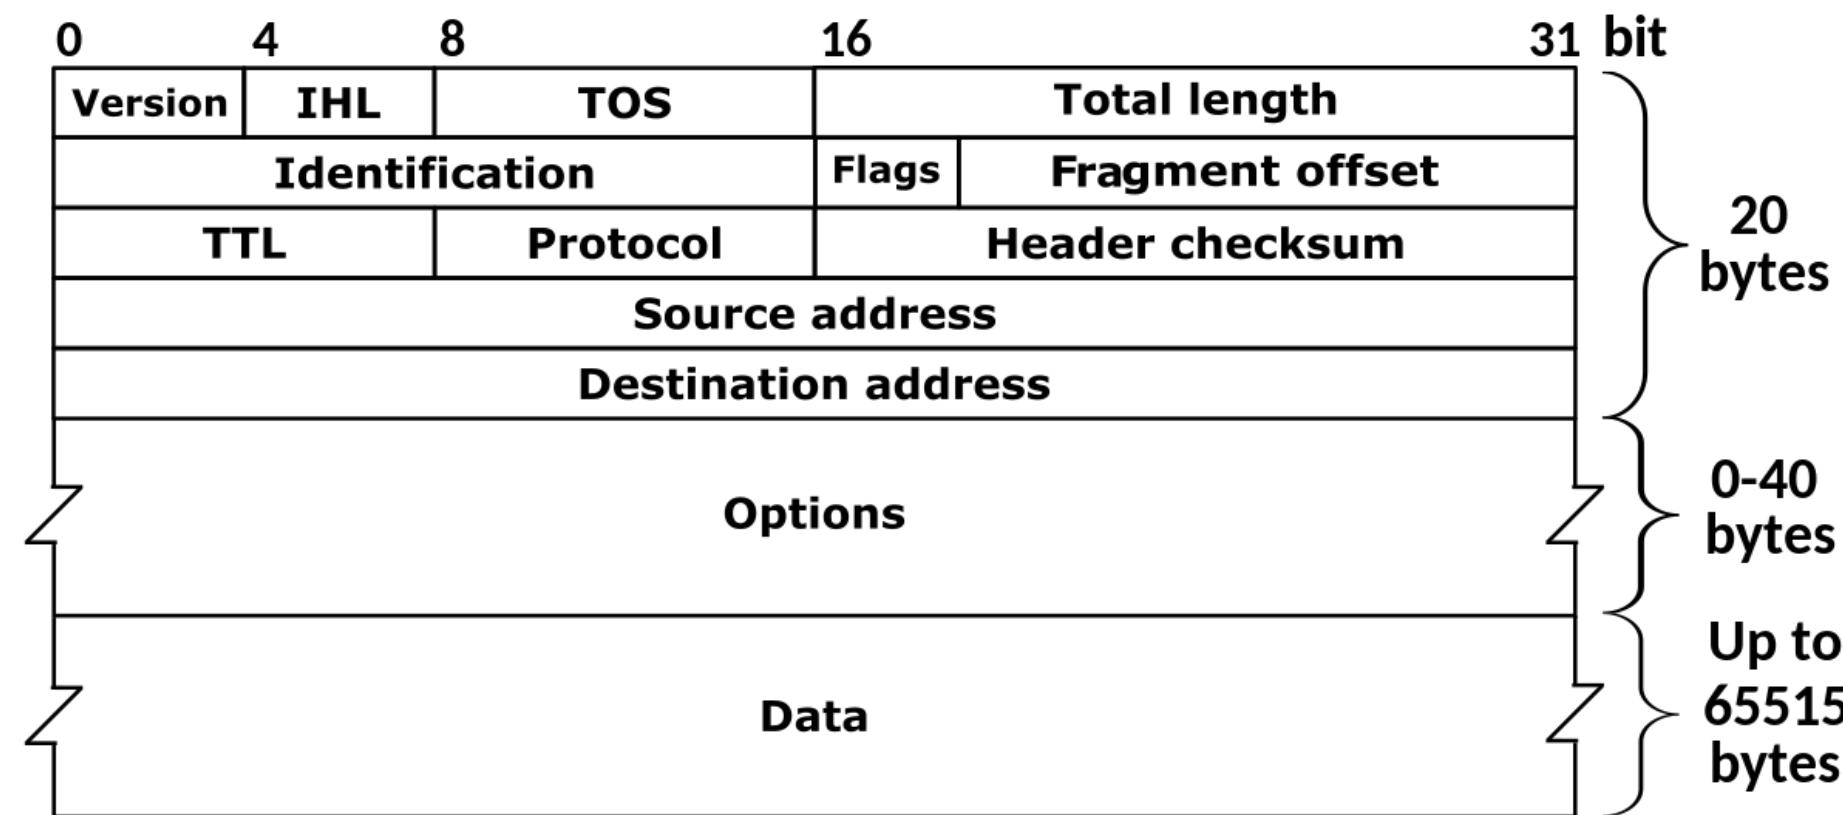
In the rest of this talk, we'll formalize my claim by building a *core calculus* of structured streaming

Windows are Streams within Streams



```
type windowed a = stream (stream a)
```

Structured Stream Elements Are Also Streams



type IPv4 = Header · Data

type Header = (version : Bit⁴) · (ihl : Bit⁴) · (dcsp : Bit⁶) ...

type Data = stream Bit

What Types Are These?

 s^*

Zero or more streams
of type S, one after
another

 $s \cdot t$

A stream of type S,
followed by stream of
type T

(Ordered Product Type)

 $\text{sing}(\tau)$

A stream which
includes exactly one
element of base type

Examples

$$(\text{sing}(\text{Int}))^*$$

Traditional Stream of ints, one after another after another.

$$((\text{sing}(\text{Int}))^*)^*$$

Windowed stream of ints

$$\text{Int} \cdot \text{Int}^*$$

Nonempty Stream of Ints

$$\text{Bit}^* \cdot \text{Bit}^*$$

Stream of bits, followed by another stream of bits

$$(\text{Bit} \cdot \text{Bit})^*$$

Stream of pairs of bits, one after another.

$$\text{Int}^* \cdot (\text{Char}^*)^* \cdot \text{Bit}$$

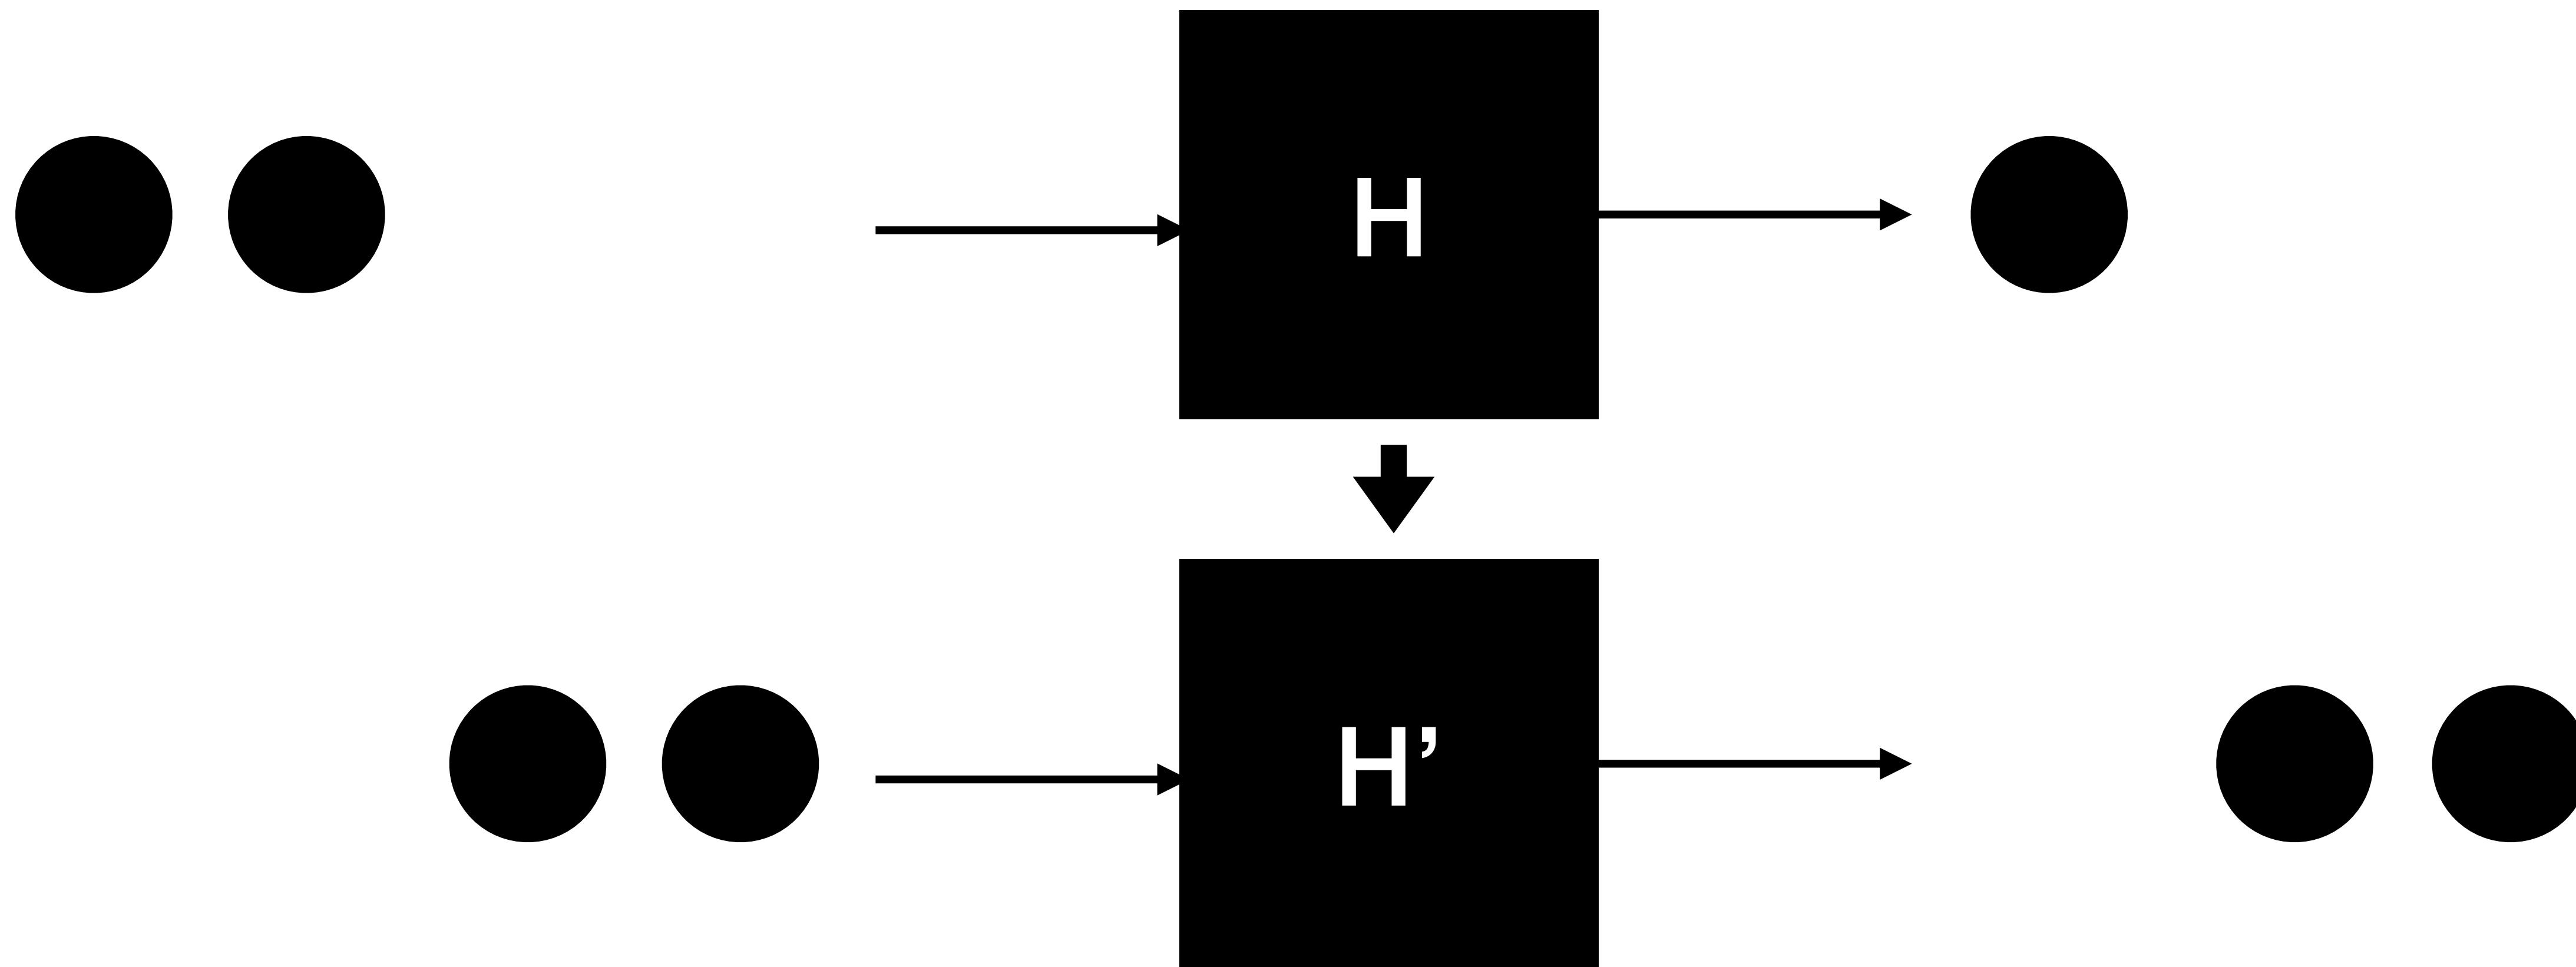
Streaming Programs Should Be Typed



Typed stream transformers must be programmable in a language with a stream-type system

$$x : S \vdash e : T$$

Transformers should have an *incremental* semantics



Typing Judgment

$$\Gamma \vdash e : S$$

*Ordered list of variables,
in sequence*

Transformer

Output Type

$$\Gamma = (x_0 : S_0); (x_1 : S_1); \dots; (x_n : S_n)$$

The diagram illustrates the structure of a stream environment Γ . It consists of a sequence of variable declarations separated by semicolons. Above the sequence, there is a horizontal line with two vertical tick marks. The left tick mark is labeled "First variable to arrive" and is positioned under the first declaration $(x_0 : S_0)$. The right tick mark is labeled "Last variable to arrive" and is positioned under the last declaration $(x_n : S_n)$.

Variables range over *structured extents* of the stream

Variable Rule

$$\Gamma; x : s; \Gamma' \vdash x : s$$

Concat Right Rule, Take 1

$$\frac{\Gamma \vdash e : S \quad \Gamma \vdash e' : T}{\Gamma \vdash (e; e') : S \cdot T}$$

Concat Right Rule, Take 1



$$\frac{\Gamma \vdash e : S \quad \Gamma \vdash e' : T}{\Gamma \vdash (e; e') : S \cdot T}$$

Concat Right Rule

$$\frac{\Gamma \vdash e : S \quad \Delta \vdash e' : T}{\Gamma ; \Delta \vdash (e; e') : S \cdot T}$$

Concat Left Rule

$$\frac{\begin{array}{c} \text{Variables for individual sub-streams} \\[10pt] \overline{\Gamma; (x : S); (y : T); \Gamma' \vdash e : R} \\[10pt] \hline \Gamma; z : S \cdot T; \Gamma' \vdash \text{let } (x; y) = z \text{ in } e : R \end{array}}{\Gamma; z : S \cdot T; \Gamma' \vdash \text{let } (x; y) = z \text{ in } e : R}$$

Variables for individual sub-streams

Variable for composite of sub-streams

Star Right Rules

$$\frac{}{\Gamma \vdash \text{nil} : S^\star}$$

$$\frac{\Gamma \vdash e : S \quad \Delta \vdash e' : S^\star}{\Gamma ; \Delta \vdash e :: e' : S^\star}$$

Star Left Rule

$$\frac{\Gamma; \Gamma' \vdash e : T \quad \Gamma; (y : S); (ys : S^*) ; \Gamma' \vdash e' : T}{\Gamma; (xs : S^*) ; \Gamma' \vdash \text{case}(xs, e, y.ys.e') : T}$$

Cut Rule

$$\frac{\Delta \vdash e : s \quad \Gamma; x : s; \Gamma' \vdash e' : t}{\Gamma; \Delta; \Gamma' \vdash \text{let } x = e \text{ in } e' : t}$$

**Does this give us the idioms of
structured streaming?**

Some Programs!

Map

```
map (f : s -> t) (xs : s*) : t* =  
  case xs of  
    | nil => nil  
    | y :: ys => (f y) :: map f ys
```

Looks like list map, runs like stream map

Some Programs!

Filter

```
filter (f : s -> bool) (xs : s*) : s* =  
  case xs of  
    | nil => nil  
    | y :: ys => if f y then y :: filter ys  
                  else filter ys
```

Some Programs!

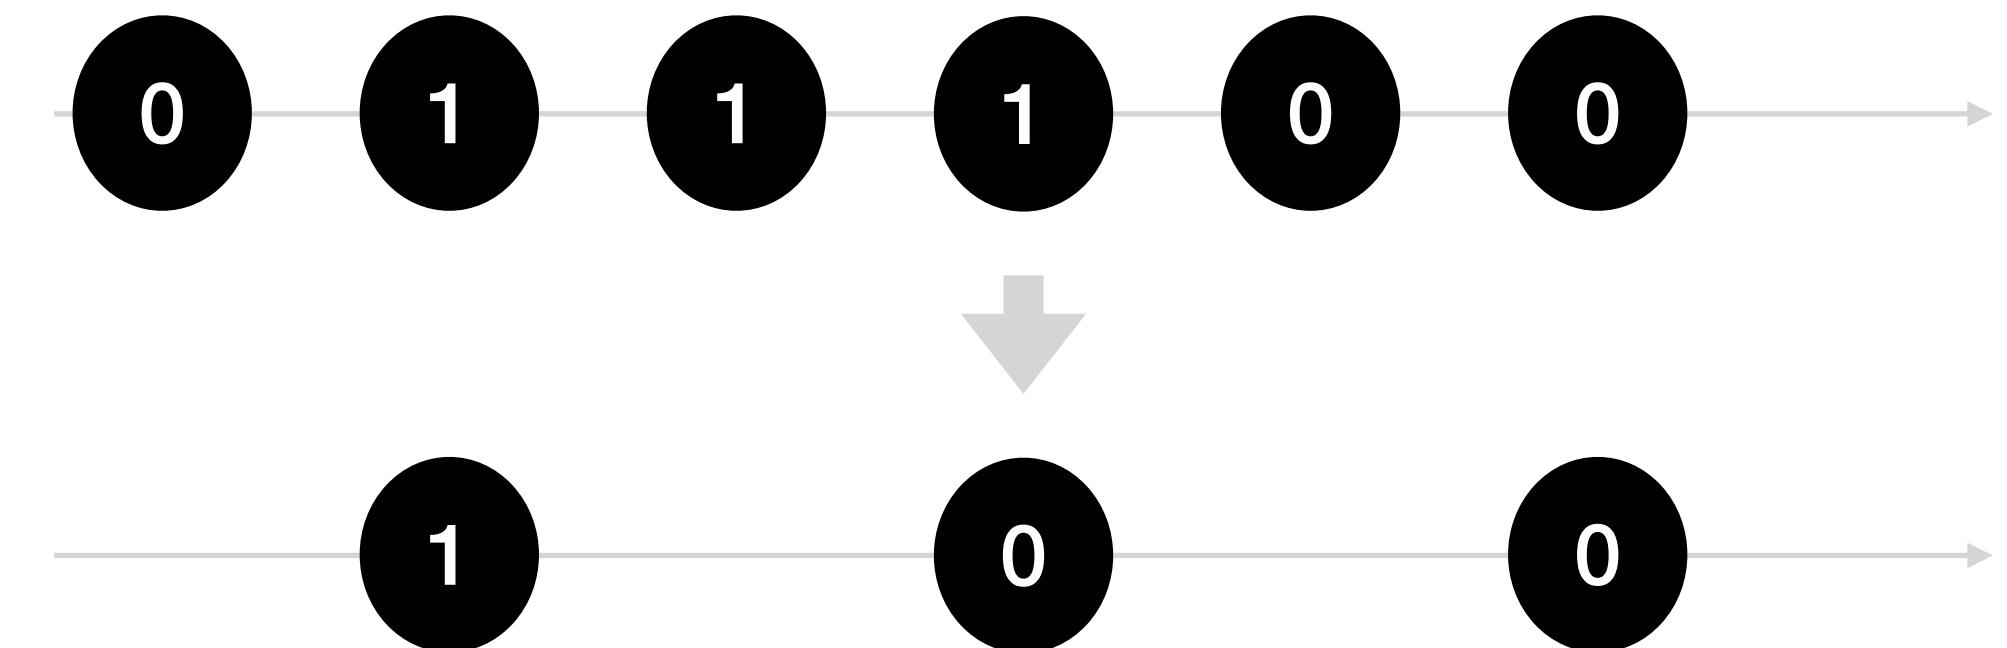
Windowed Aggregation

```
Windowed Operation           Stream of Windows
|                           |
|-----|-----|
windowAgg (f : s* -> t) (xs : s**) : t* =
map f xs
```

Windowed aggregation is a map over a stream of windows

Some Programs!

Pairwise Xor of a Stream



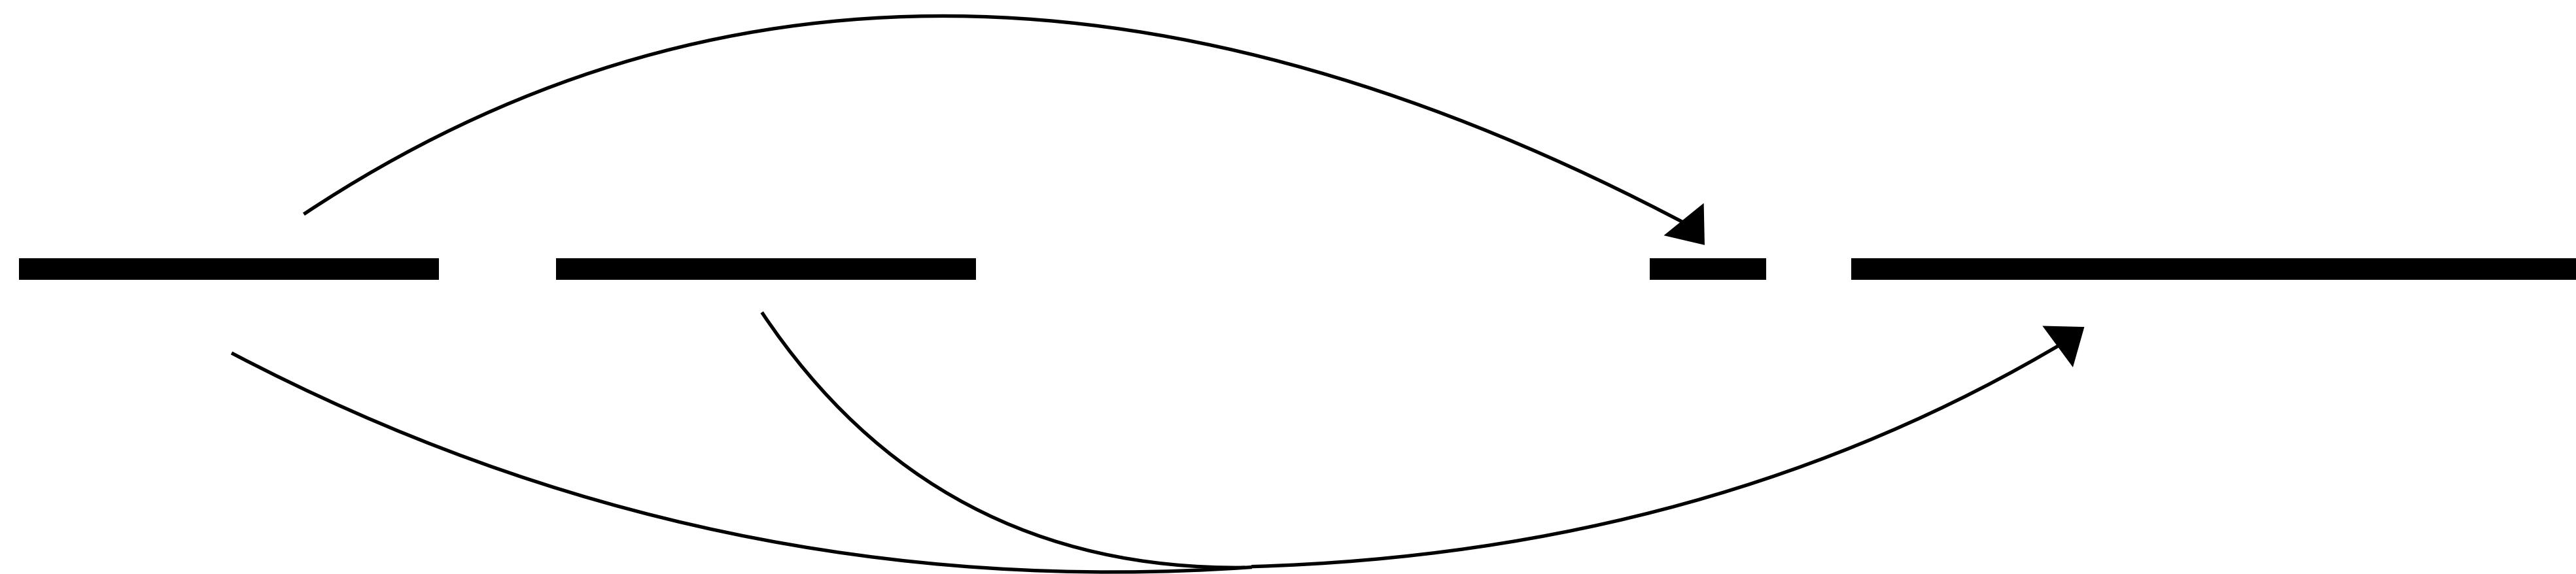
```
parse (xs : bit*) : (bit . bit)* =
case xs of
| nil => nil
| [_] => nil
| b::b'::bs => (b;b') :: parse bs

xor (x : bit . bit) : bit
let (a;b) = x in
if a then !b else b
```

```
pairssxor (xs : bit*) : bit* =
map xor (parse xs)
```

What About State?

Mental Model of *Stateful* Structured Streaming



Earlier extents of the input can be used to produce later extents of output

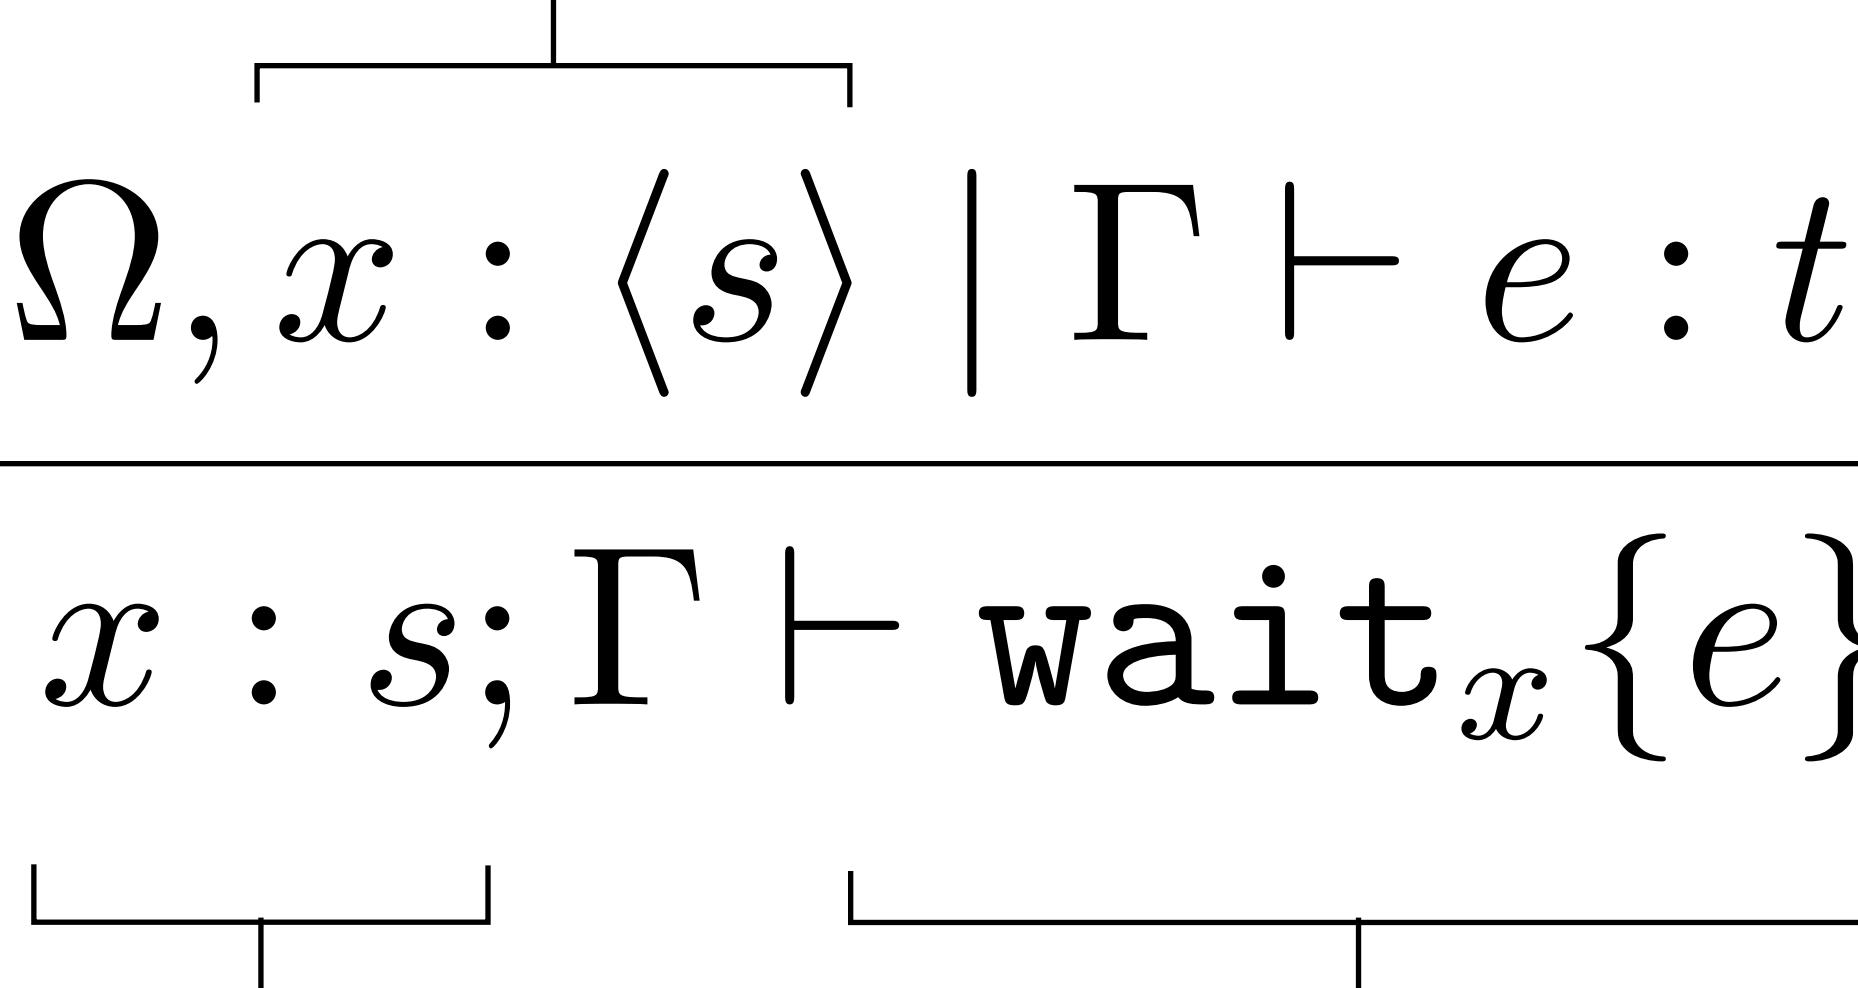
State Context

$$\Omega \mid \Gamma \vdash e : S$$

Fully structural context bound to “already-seen data”

Wait Rule

$$\frac{\begin{array}{c} x \text{ now in the past} \\[10pt] \overline{\quad\quad\quad} \\[10pt] \Omega, x : \langle s \rangle \mid \Gamma \vdash e : t \end{array}}{\Omega \mid x : s ; \Gamma \vdash \text{wait}_x\{e\} : t}$$



Stream with
substream bound to x Wait for x then run e

Using State Variables

$$\Omega, x : \langle S \rangle \mid \Gamma \vdash x : S$$

Examples, With State!

Some Programs!

Window Creation (Sliding)

```
slidingWindow {prev : s*} (xs : s*) : s** =
  case xs of
    | nil => nil
    | y :: ys => wait y {
        let next = update prev y in
        next :: slidingWindow next ys
      }
```

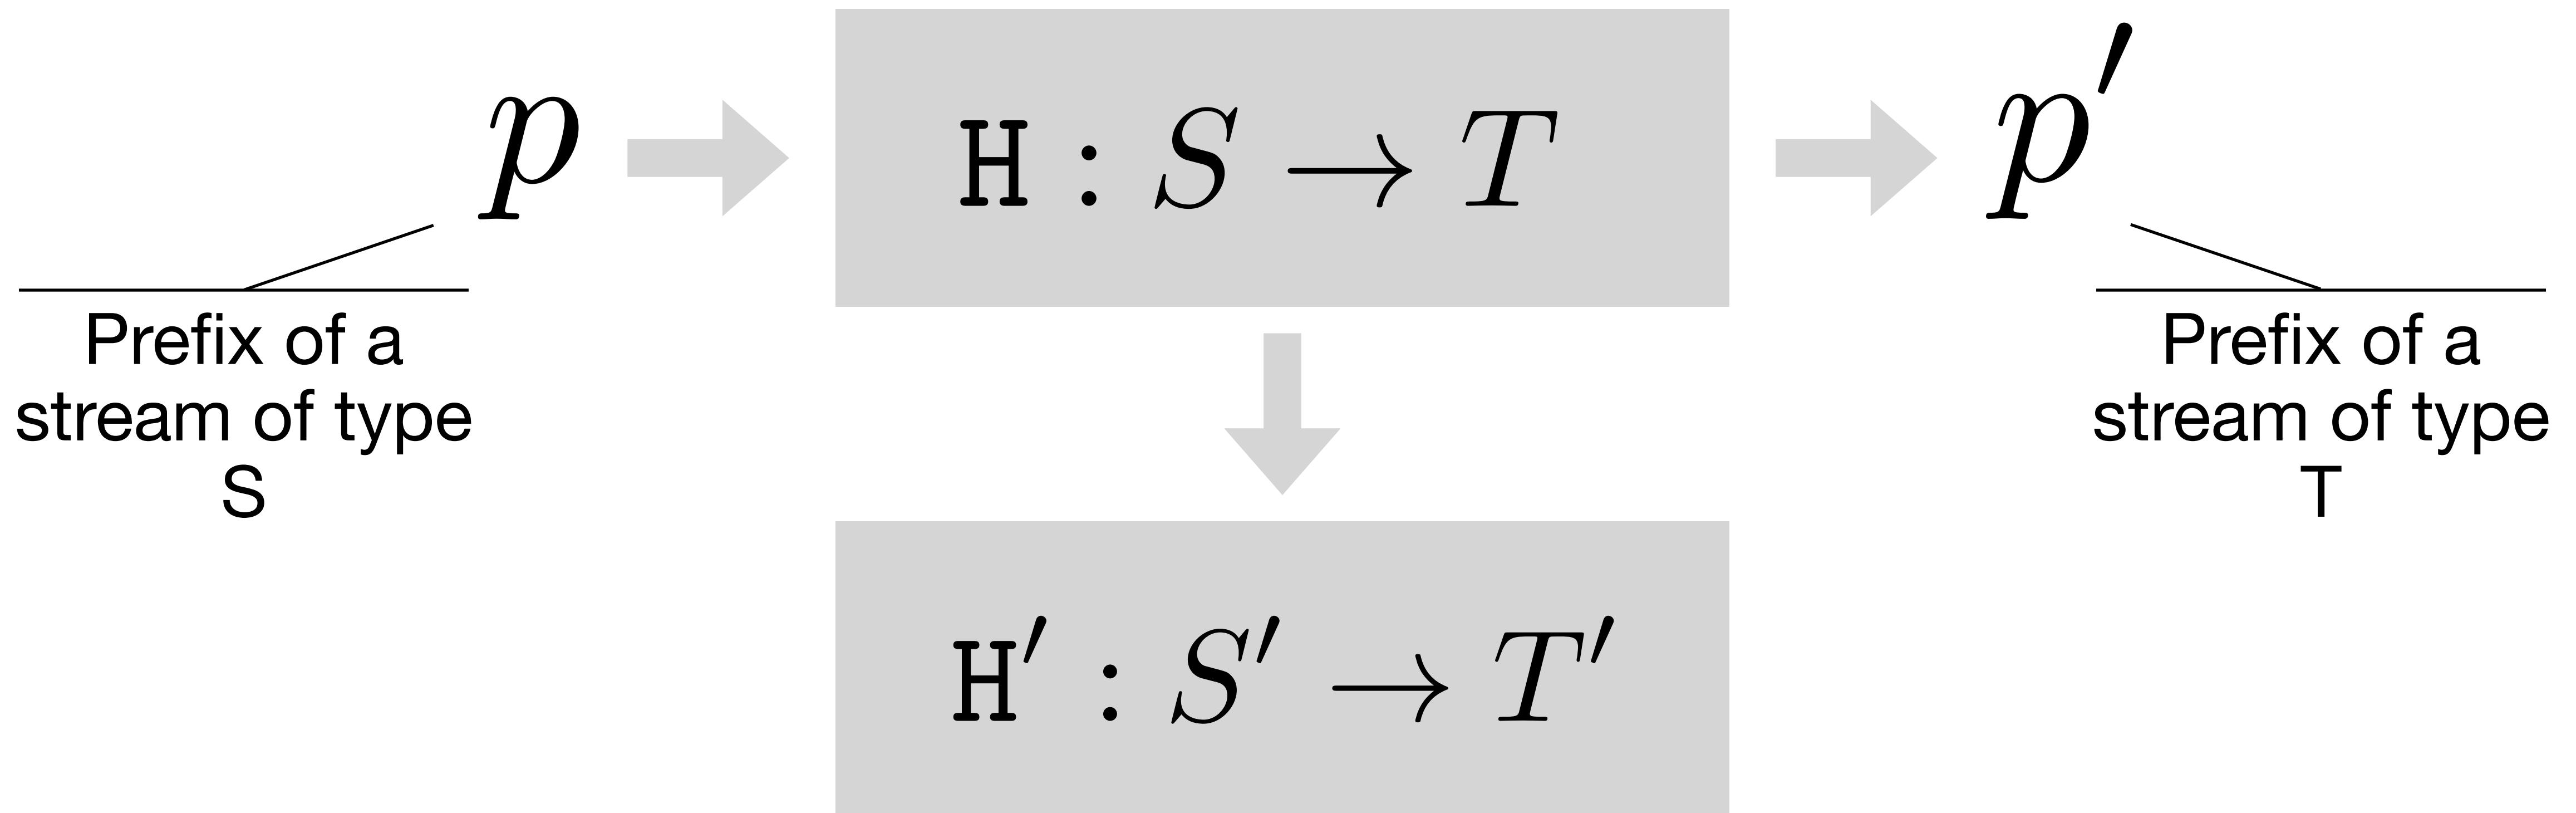
Some Programs!

Window Creation (Tumbling)

```
tumblingWindow {k : int} {acc : s*} (xs : s*) : s** =
  case xs of
    | nil => nil
    | y :: ys => if acc.length == k then
      (y :: acc) :: tumblingWindow k [] ys
    else
      wait y {
        tumblingWindow k (y :: acc) ys
      }
```

Semantics

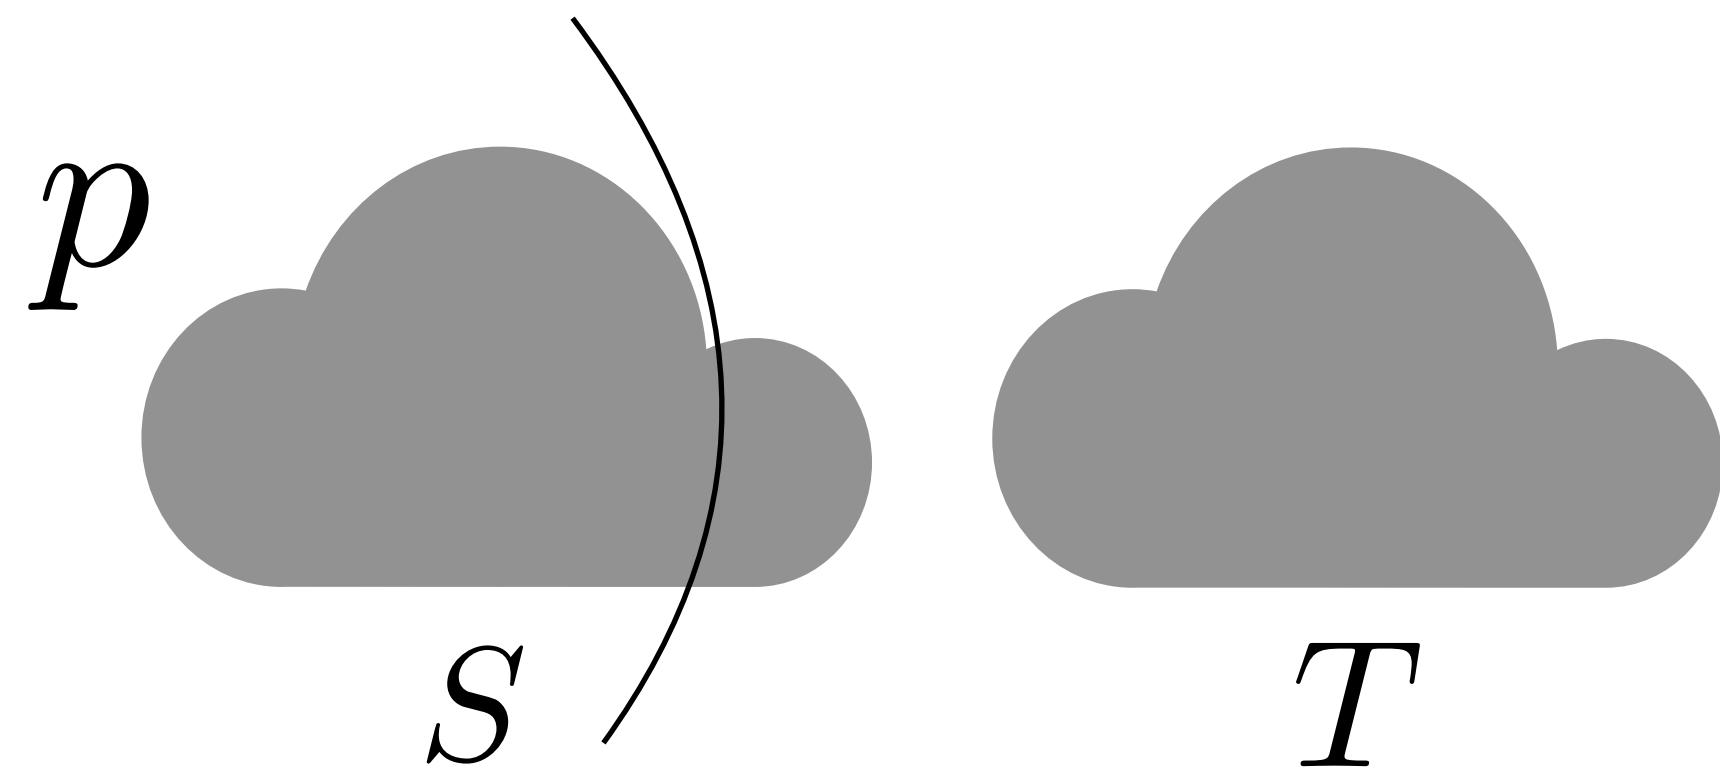
Semantics



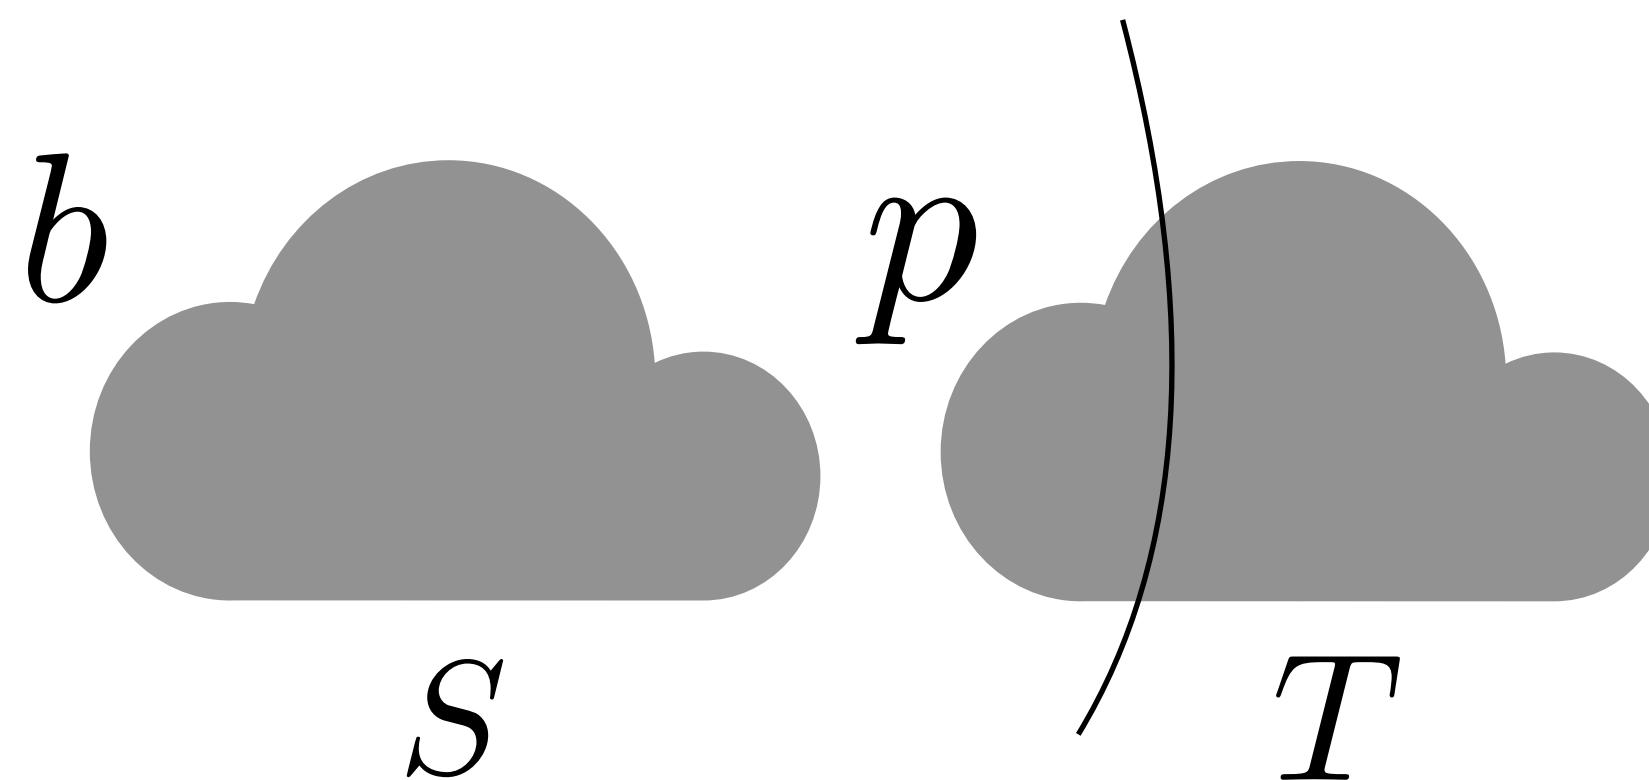
New transformer, accepting the rest of S , and producing the rest of T

What is a Prefix of a Stream of Type S?

Concat



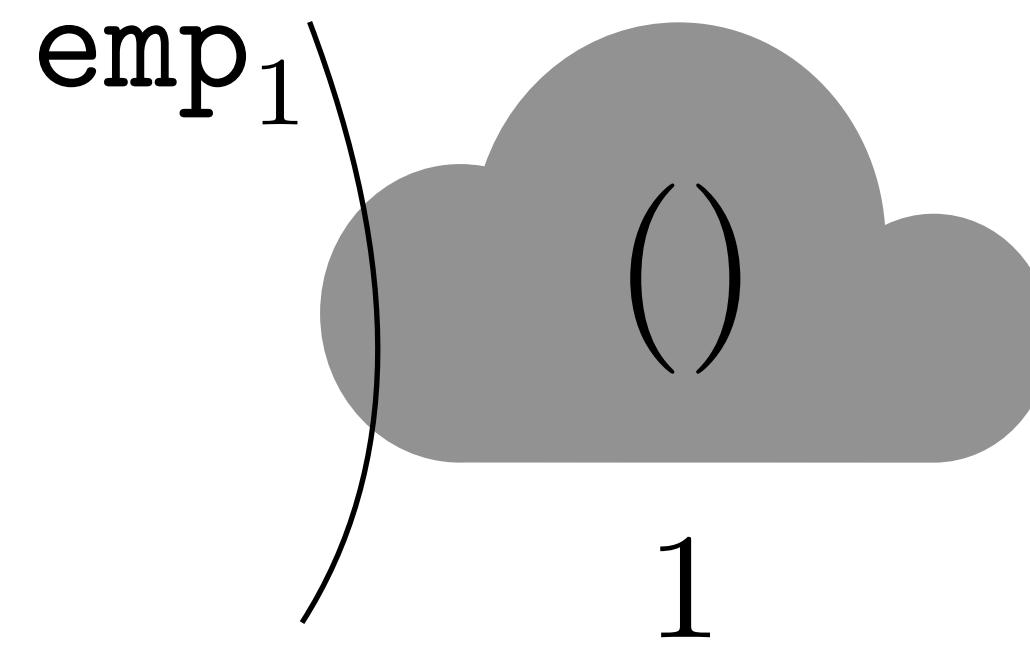
$$\frac{p : \text{prefix}(S)}{\text{inl}(p) : \text{prefix}(S \cdot T)}$$



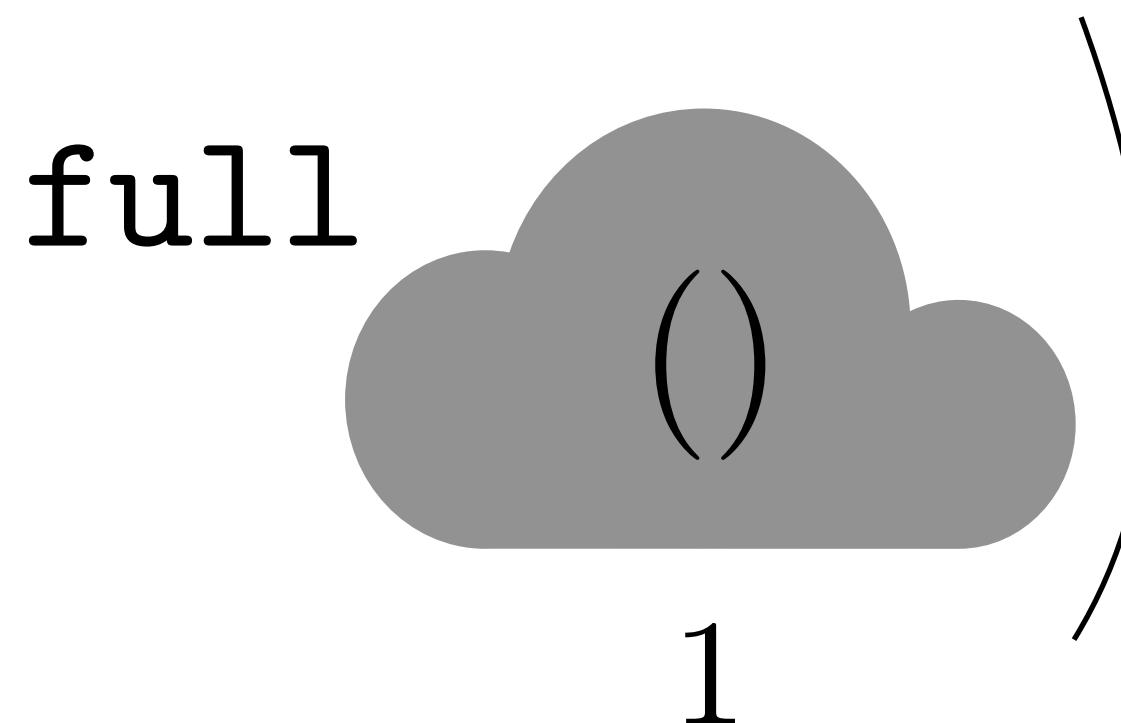
$$\frac{b : \text{batch}(S) \quad p : \text{prefix}(T)}{\text{inr}(b, p) : \text{prefix}(S \cdot T)}$$

What is a Prefix of a Stream of Type S?

Base Type (1)

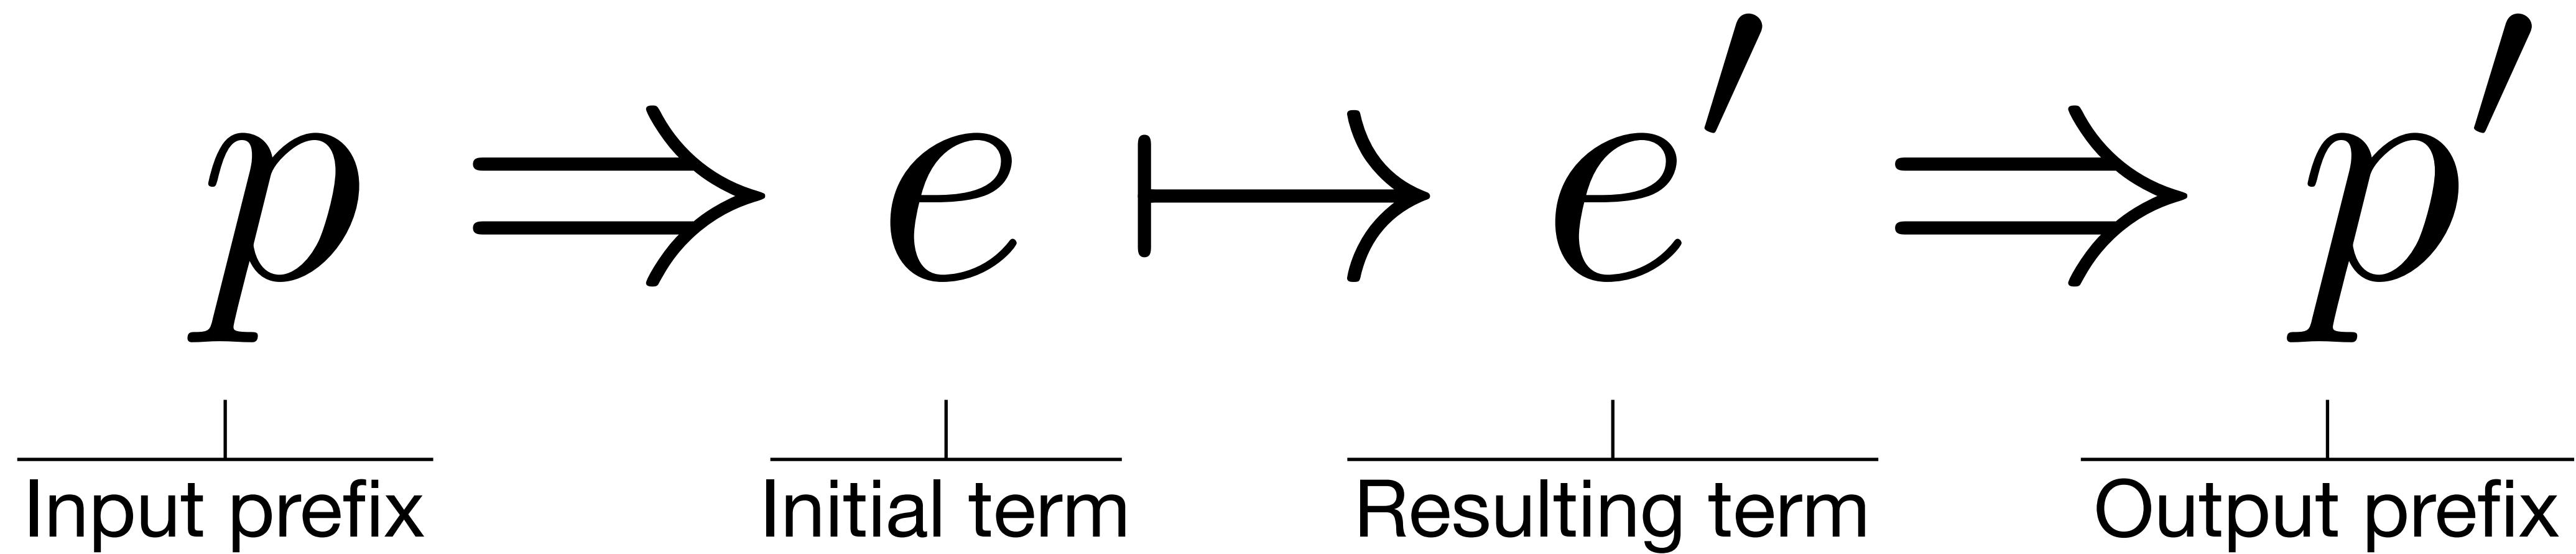


$\text{emp}_1 : \text{prefix}(1)$



$\text{full} : \text{prefix}(1)$

Semantics



Preservation

If:

$$\left[\begin{array}{l} \cdot \mid \Gamma \vdash e : S \\ p : \text{prefix}(\Gamma) \end{array} \right]$$

$$p \Rightarrow e \mapsto e' \Rightarrow p'$$

$$p' : \text{prefix}(S)$$

Then:

$$\left[\delta_p \Gamma \vdash e' : \delta_{p'} S \right]$$

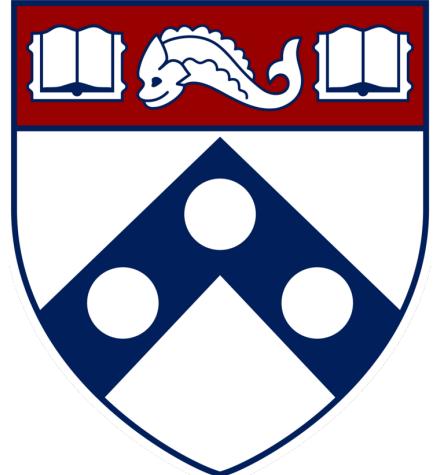
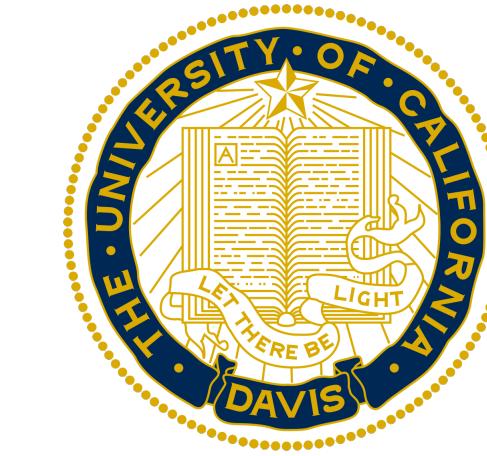
Progress

If:

$$\left[\begin{array}{l} \cdot \mid \Gamma \vdash e : S \\ p : \text{prefix}(\Gamma) \end{array} \right]$$

Then: $\exists p, e'. p \Rightarrow e \mapsto e' \Rightarrow p'$

Thank you!



Type System

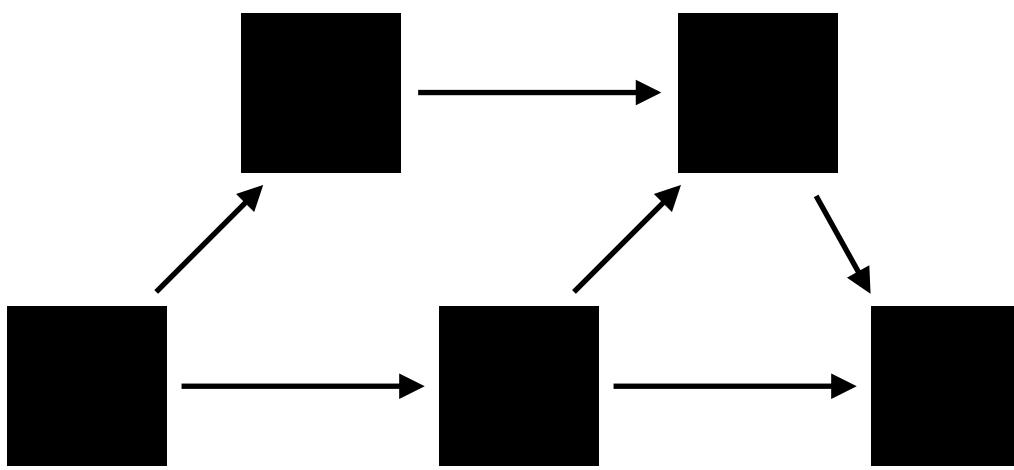
$$\Gamma \vdash e : S$$

Operational Semantics

$$\eta \Rightarrow e \mapsto e' \Rightarrow p$$

Expressiveness

*CQL ->
Stream Types*



Low-Level Dataflow Model

Mechanized Metatheory

$$\begin{aligned} \Gamma \vdash e : S \\ p : \text{prefix}(\Gamma) \end{aligned}$$

.....

$$\exists p, e'. p \Rightarrow e \mapsto e' \Rightarrow p'$$

Interpreter in Haskell

```
{-  
G |- e1 : s  D |- e2 : t  
-----  
G;D |- (e1;e2) : s . t  
-}  
eval (CatPfxA p) (CatPair e1 e2) = do  
  ER {output = p', newExpr = e1'} ← eval p e1  
  return $ ER {output = CatPfxA p', newExpr = CatPair e1' e2}  
  
eval (CatPfxB b p) (CatPair e1 e2) = do  
  b' ← batchEval b e1  
  ER {output = p', newExpr = e2'} ← eval p e2  
  return $ ER {output = CatPfxB b' p', newExpr = e2'}
```